

Cafe Sales Prediction

Predicting cafe sales using historical transaction data.

In [304]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

plt.style.use('dark_background')
df = pd.read_csv('/content/dirty_cafe_sales.csv')
df.head(100)
```

Out [304]:

| | Transaction ID | Item | Quantity | Price Per Unit | Total Spent | Payment Method | Location | Transaction Date |
|-----|----------------|--------|----------|----------------|-------------|----------------|----------|------------------|
| 0 | TXN_1961373 | Coffee | 2 | 2.0 | 4.0 | Credit Card | Takeaway | 2023-09-08 |
| 1 | TXN_4977031 | Cake | 4 | 3.0 | 12.0 | Cash | In-store | 2023-05-16 |
| 2 | TXN_4271903 | Cookie | 4 | 1.0 | ERROR | Credit Card | In-store | 2023-07-19 |
| 3 | TXN_7034554 | Salad | 2 | 5.0 | 10.0 | UNKNOWN | UNKNOWN | 2023-04-27 |
| 4 | TXN_3160411 | Coffee | 2 | 2.0 | 4.0 | Digital Wallet | In-store | 2023-06-11 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 95 | TXN_8268061 | Salad | 3 | 5.0 | 15.0 | ERROR | Takeaway | 2023-08-20 |
| 96 | TXN_5220895 | Salad | 5 | 5.0 | 25.0 | Cash | In-store | 2023-06-10 |
| 97 | TXN_3085509 | Coffee | 4 | 2.0 | 8.0 | Digital Wallet | In-store | 2023-04-15 |
| 98 | TXN_9999113 | Juice | 4 | 3.0 | 12.0 | Cash | Takeaway | 2023-05-27 |
| 99 | TXN_8779771 | Coffee | 4 | 2.0 | 8.0 | Cash | In-store | 2023-07-25 |

100 rows × 8 columns

Data Preprocessing and EDA

In [305]:

```
import numpy as np
df.replace(['ERROR', 'UNKNOWN'], np.nan, inplace=True)
```

In [306]:

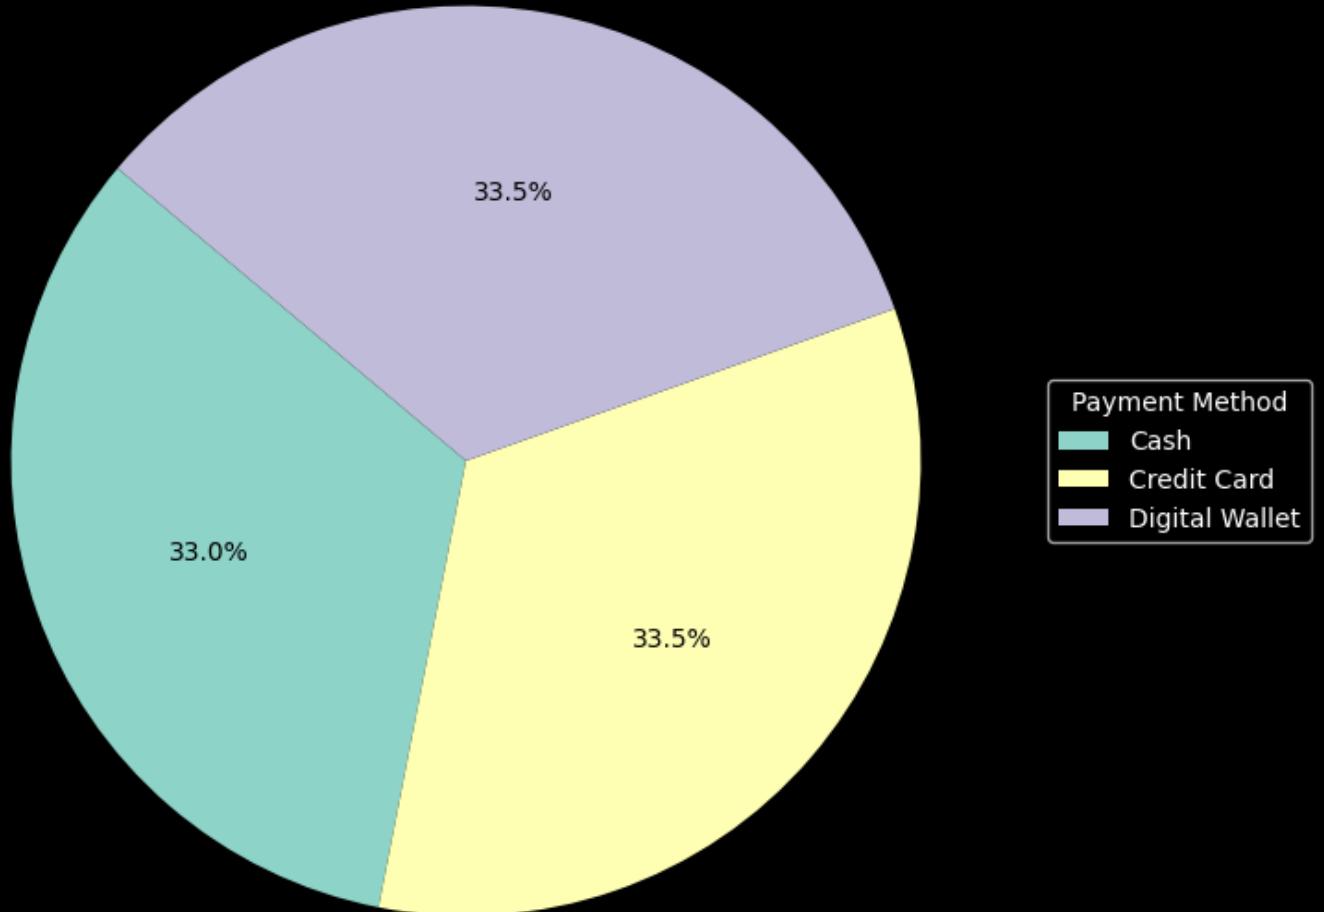
```
df.columns = df.columns.str.replace(' ', '_')
```

In [307]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Payment_Method', 'Total_Spent'], inplace=True)
payment_method_totals = df_copy.groupby('Payment_Method')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(payment_method_totals, labels=payment_method_totals.index,
                                    autopct='%.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Payment Method')
plt.legend(wedges, payment_method_totals.index, title="Payment Method", loc="center left",
           bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```



Total Spent by Payment Method

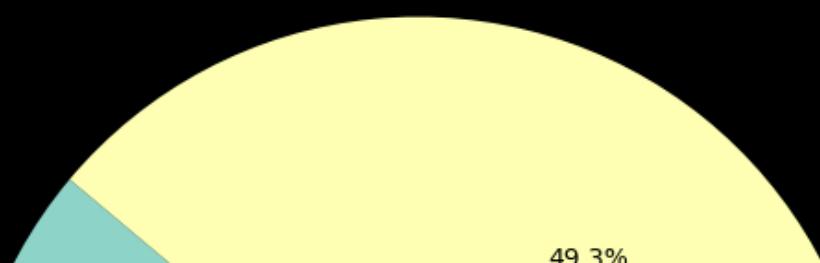


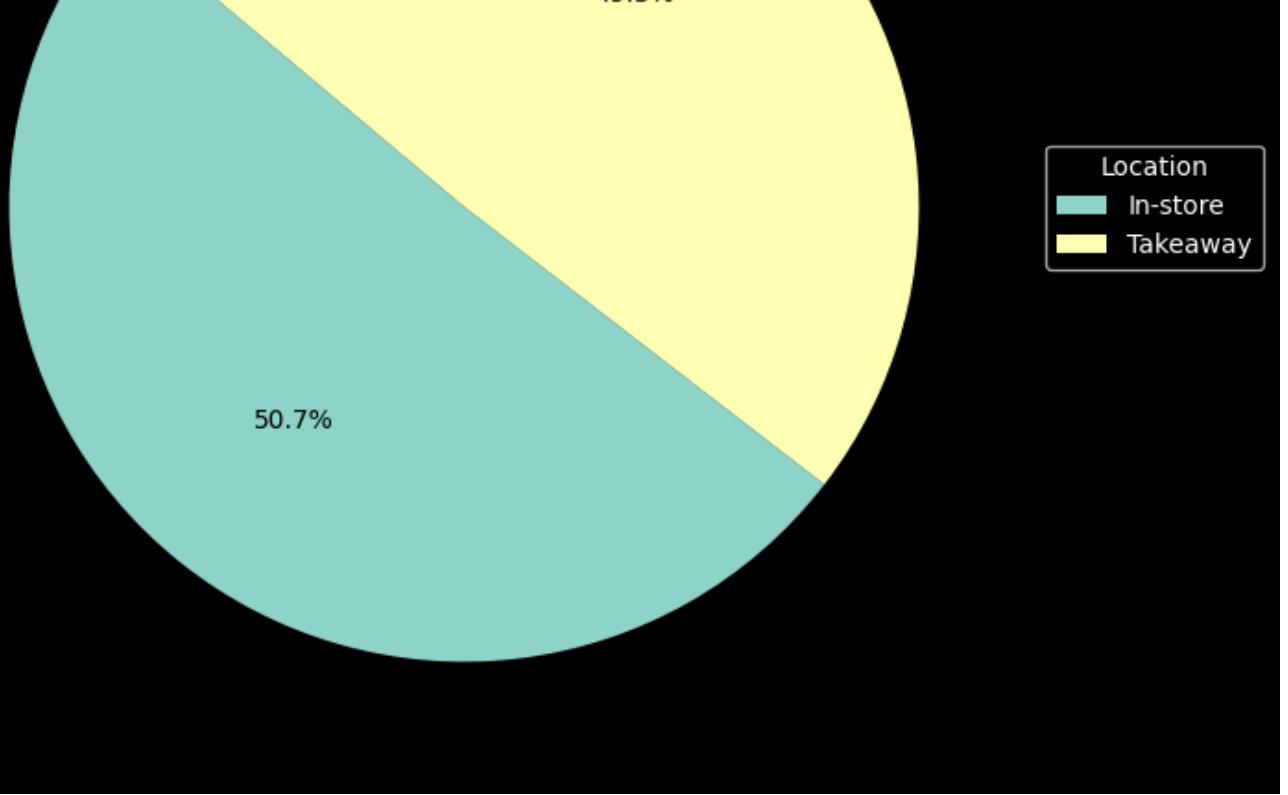
After analyzing the distribution of the Payment Method feature, we observe that the total spending is almost evenly divided among the three methods (approximately 33%, 33.5%, and 33.5%). This indicates that the Payment Method does not significantly impact the target variable (Total Spent), and therefore, it can be considered for removal as a non-informative feature.

In [308]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Location', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Location')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Location')
plt.legend(wedges, location_totals.index, title="Location", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```

Total Spent by Location

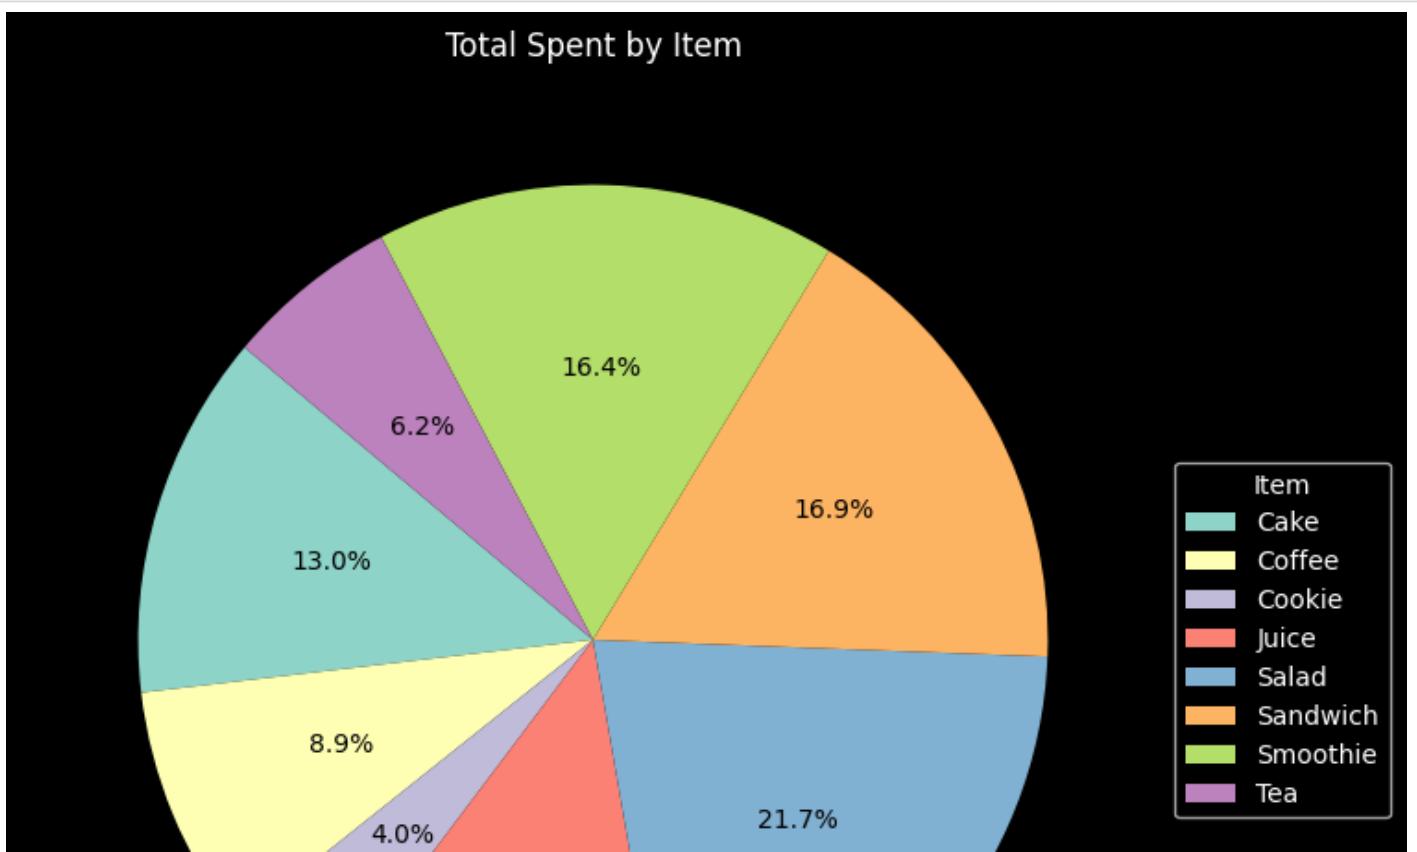


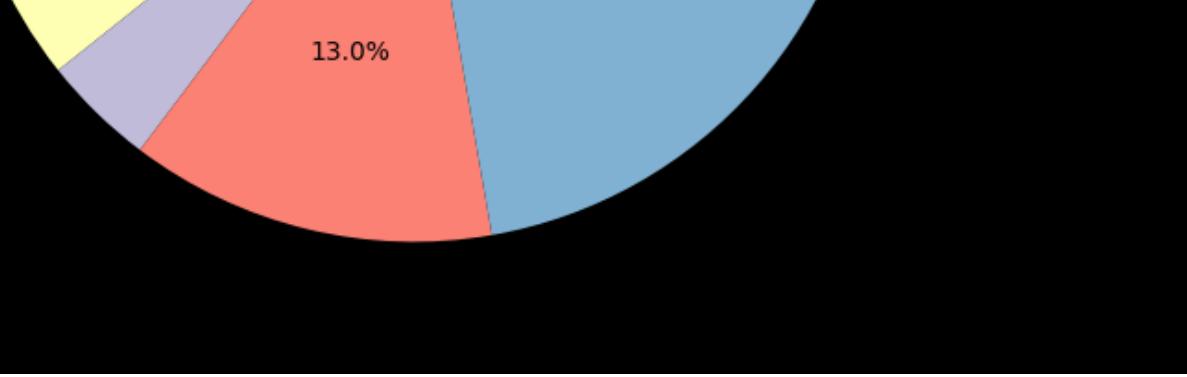


The analysis of the Location feature shows that total spending is nearly evenly split between In-Store (50.7%) and Takeaway (49.3%). This minimal difference suggests that Location has little to no impact on the target variable (Total Spent), and thus, it can be considered for removal as an uninformative feature.

In [309]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Item', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Item')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Item')
plt.legend(wedges, location_totals.index, title="Item", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```



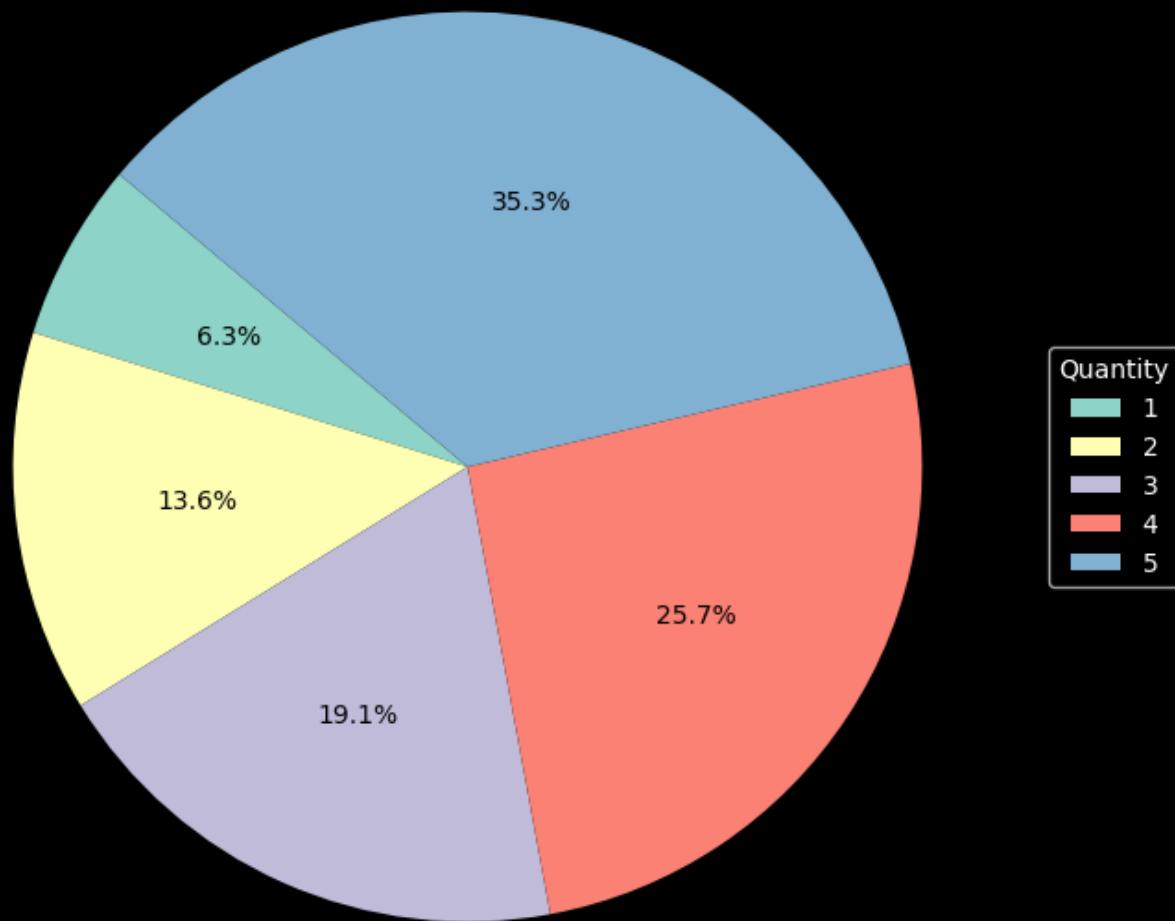


We observe that there are 8 different items, and the percentage of total spending across these items varies significantly — ranging from 4% to 21.7%. This high variability indicates that the Item feature plays an important role in our analysis and should be retained for further study.

In [310]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Quantity', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Quantity')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Quantity')
plt.legend(wedges, location_totals.index, title="Quantity", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```

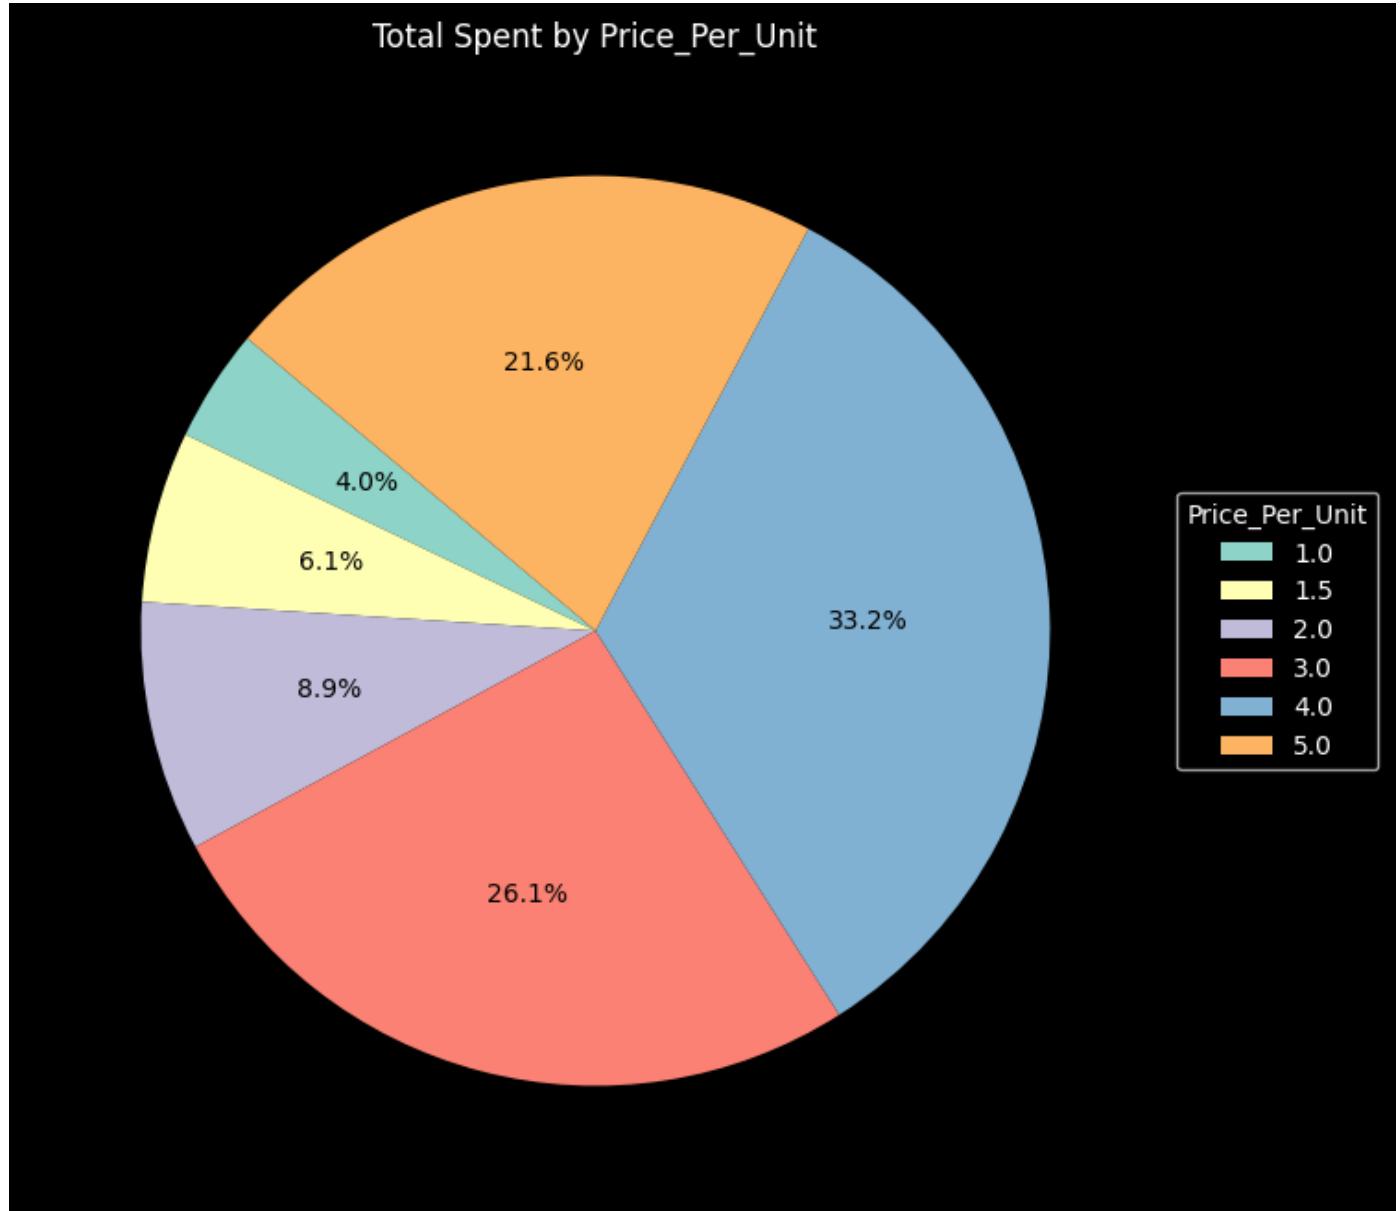
Total Spent by Quantity



The quantity feature significantly influences the target variable, total spent. When the total spent percentage is divided into five levels ranging from 1 to 5, it varies substantially—from as low as 6.3% up to 35.3%. This indicates that as the quantity increases, the total spent tends to rise noticeably, highlighting a strong positive relationship between quantity and total Spent.

In [311]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Price_Per_Unit', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Price_Per_Unit')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Price_Per_Unit')
plt.legend(wedges, location_totals.index, title="Price_Per_Unit", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```



The price per unit feature plays an important role in determining the total spent. When the total spent percentage is divided into five levels from 1 to 5, it varies widely—from 6.3% up to 35.3%. This variation shows that changes in the price per unit are closely associated with changes in total spending, indicating that higher unit prices generally lead to higher total expenditures.

In [312]:

```
nan_transaction_date_count = df['Transaction_Date'].isna().sum()
print(f"Number of rows where 'transaction_date' is NaN: {nan_transaction_date_count}")
```

Number of rows where 'transaction_date' is NaN: 460

Out of 10,000 rows, 460 have a missing transaction_date value, representing only 4.6% of the data. Since this is a relatively small proportion, these rows can be safely dropped without significantly impacting the overall analysis.

In [313]:

```
df.dropna(subset=['Transaction_Date'], inplace=True)
```

In [314]:

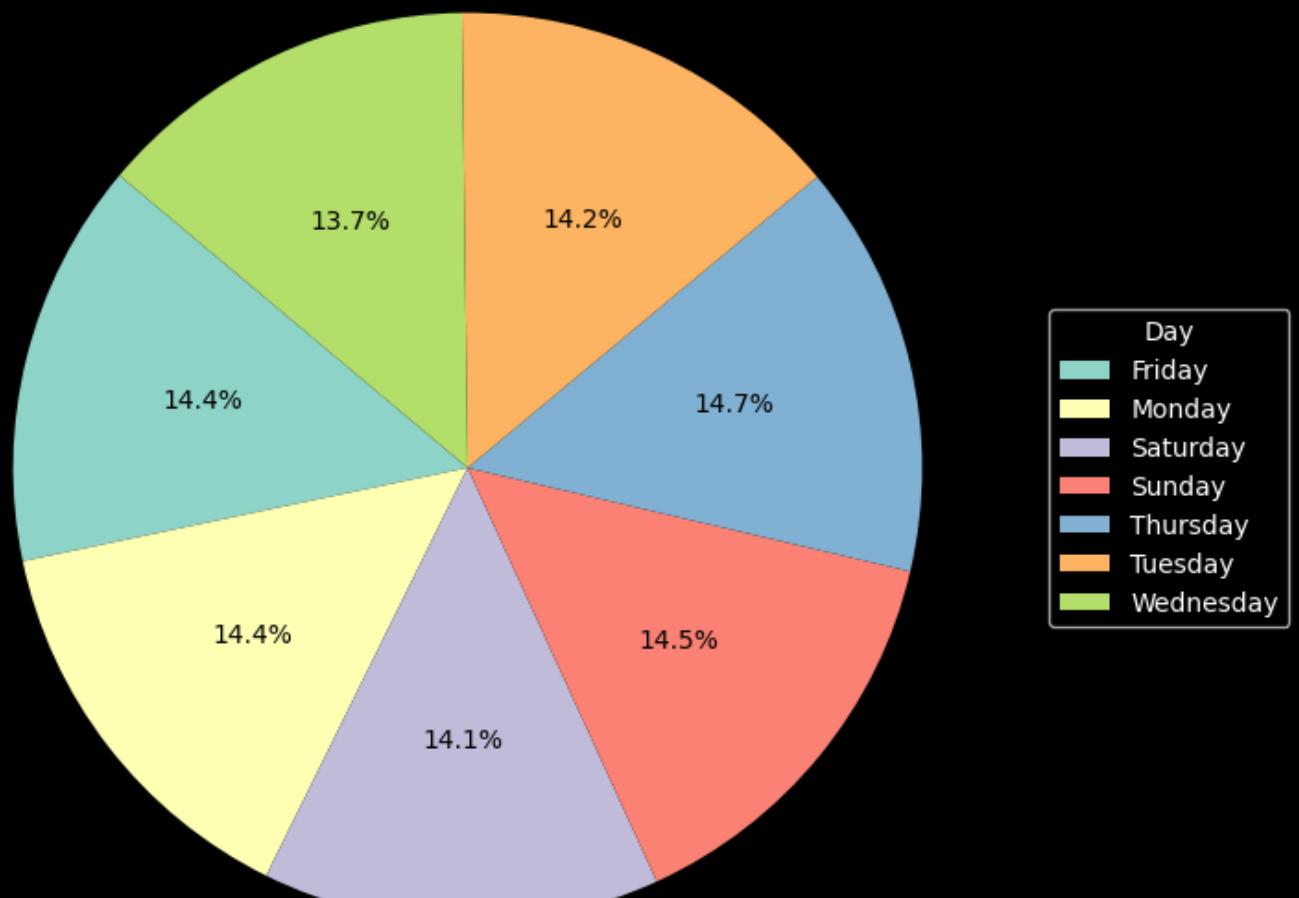
```
df['Transaction_Date'] = pd.to_datetime(df['Transaction_Date'])
df['Day'] = df['Transaction_Date'].dt.day_name()
df['Month'] = df['Transaction_Date'].dt.month
```

To enhance the analysis, I plan to extract the day of the week and month from the transaction_date feature. These time-based features are likely to be important, as customer behavior can vary depending on the day or month. For example, sales may be higher on weekends or during specific months. By creating these features, we can later visualize the data to verify whether such patterns exist.

In [315]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Day', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Day')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Day')
plt.legend(wedges, location_totals.index, title="Day", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```

Total Spent by Day



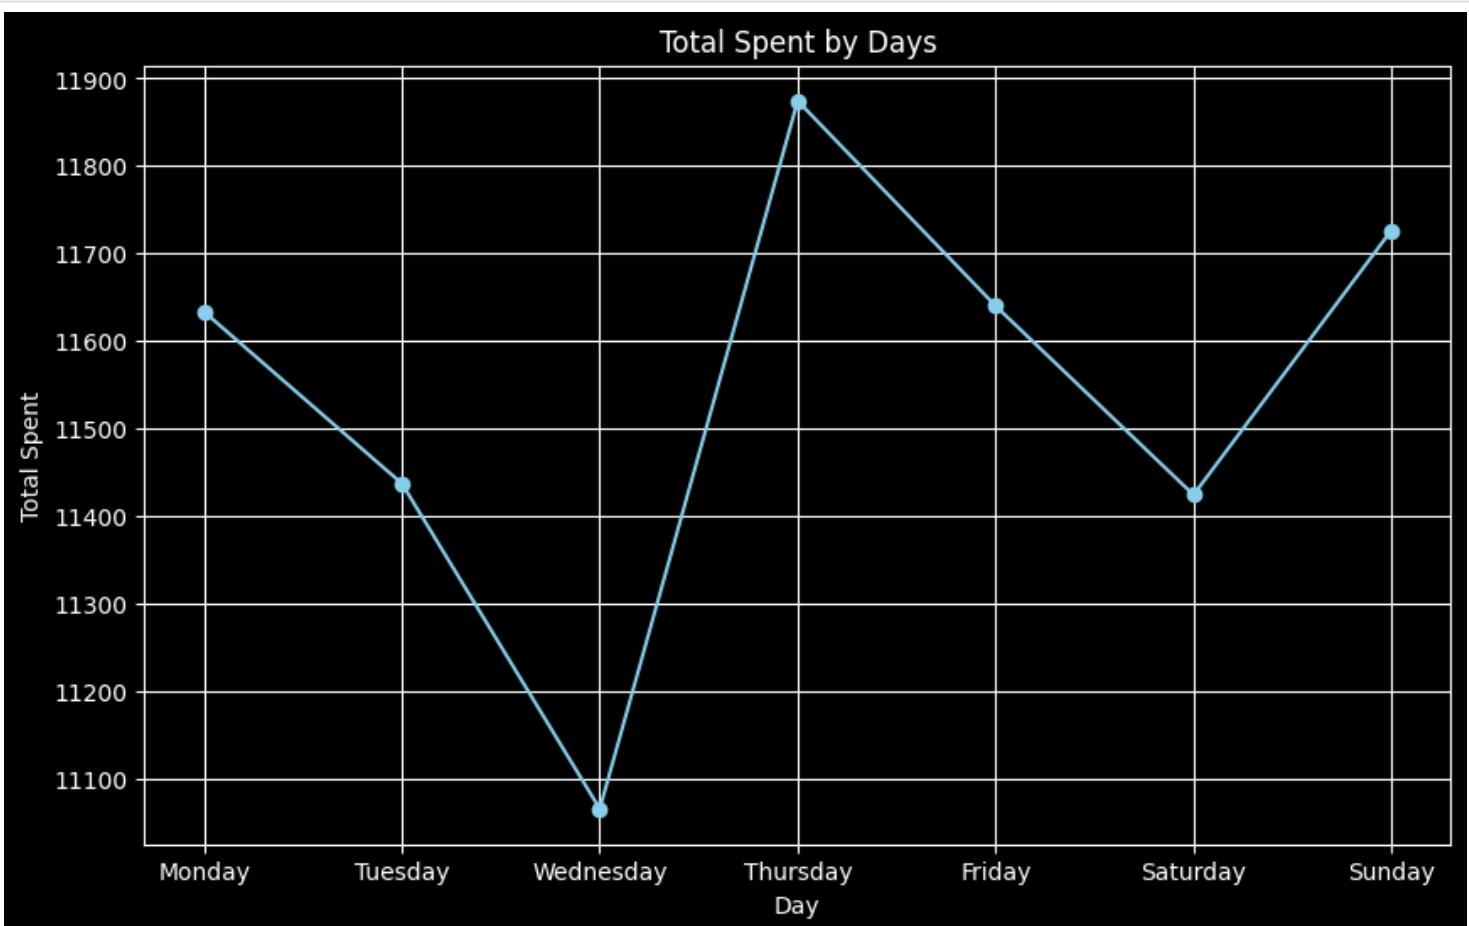
In [316]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Day', 'Total_Spent'], inplace=True)
month_totals = df_copy.groupby('Day')['Total_Spent'].sum()

# Define the desired order of the days
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']

# Reindex the month_totals Series to reflect the desired order
month_totals = month_totals.reindex(day_order)

plt.figure(figsize=(10, 6))
plt.plot(month_totals.index, month_totals.values, marker='o', linestyle='-', color='skyblue')
plt.title('Total Spent by Days')
plt.xlabel('Day')
plt.ylabel('Total Spent')
plt.xticks(month_totals.index)
plt.grid(True)
plt.show()
```

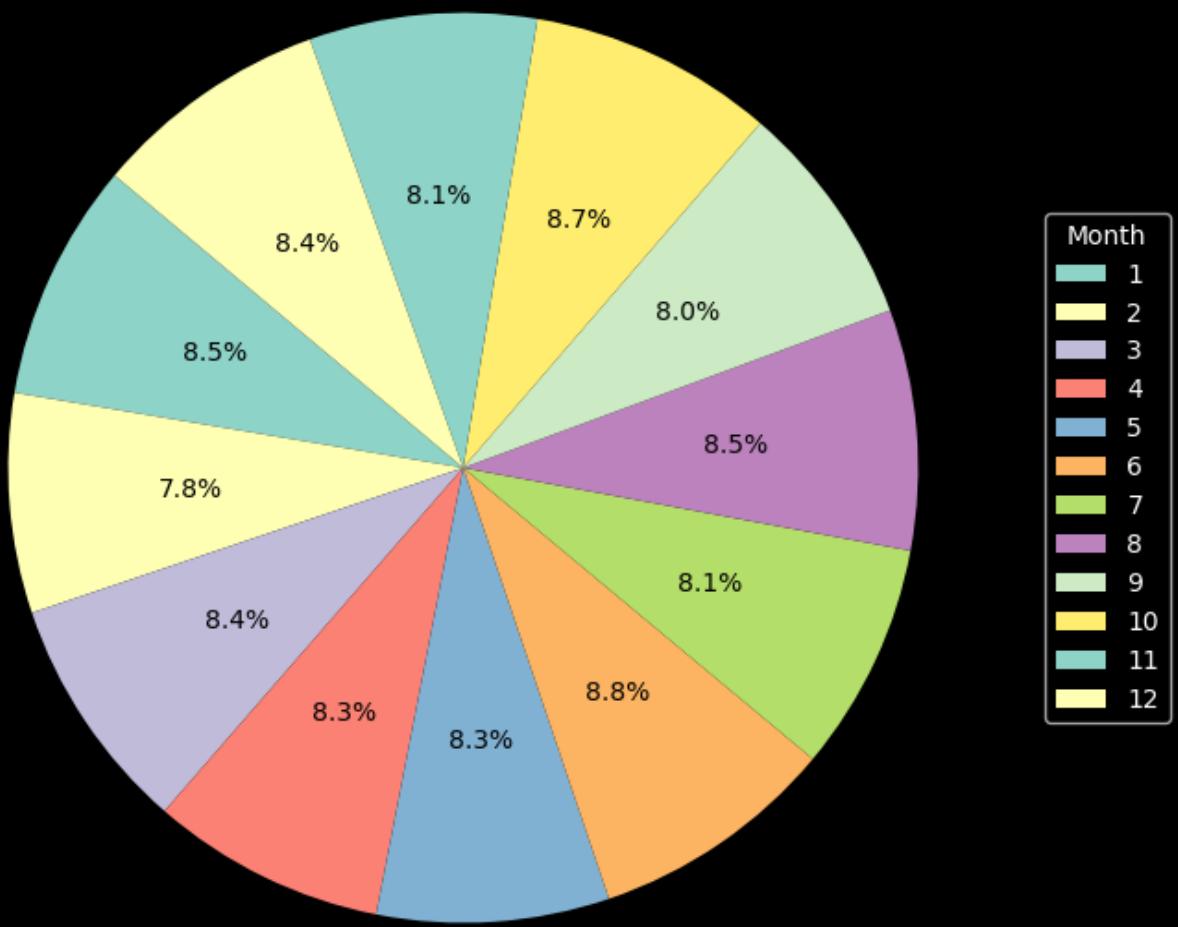


The total spent is not very affected by which day it is. But it is still considerable.

In [317]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Month', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Month')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Month')
plt.legend(wedges, location_totals.index, title="Month", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```

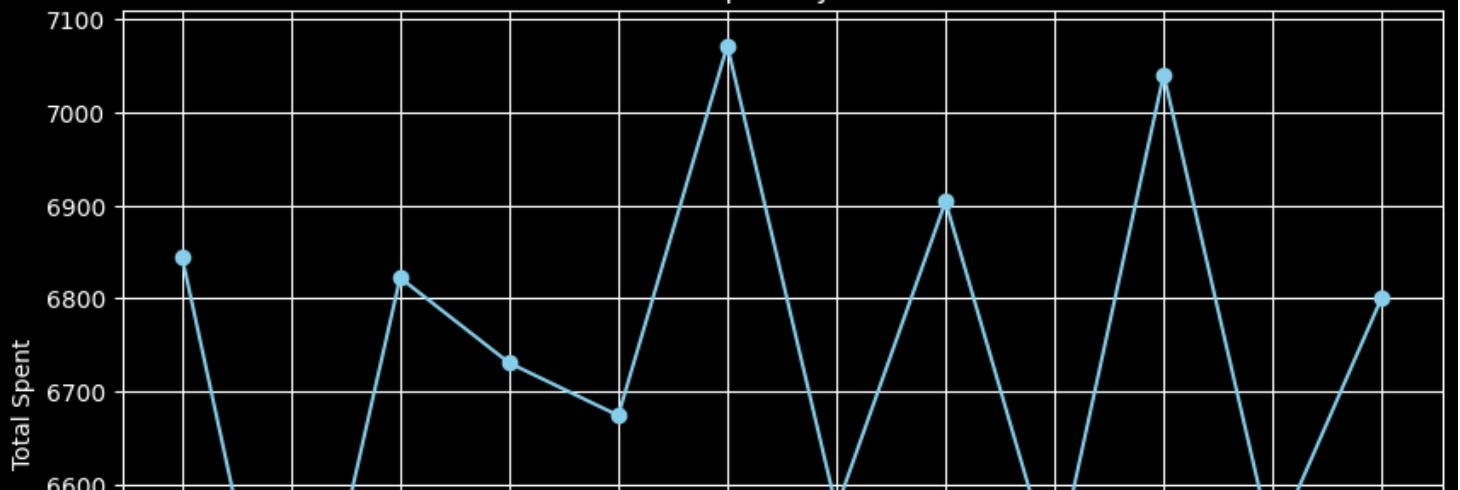
Total Spent by Month

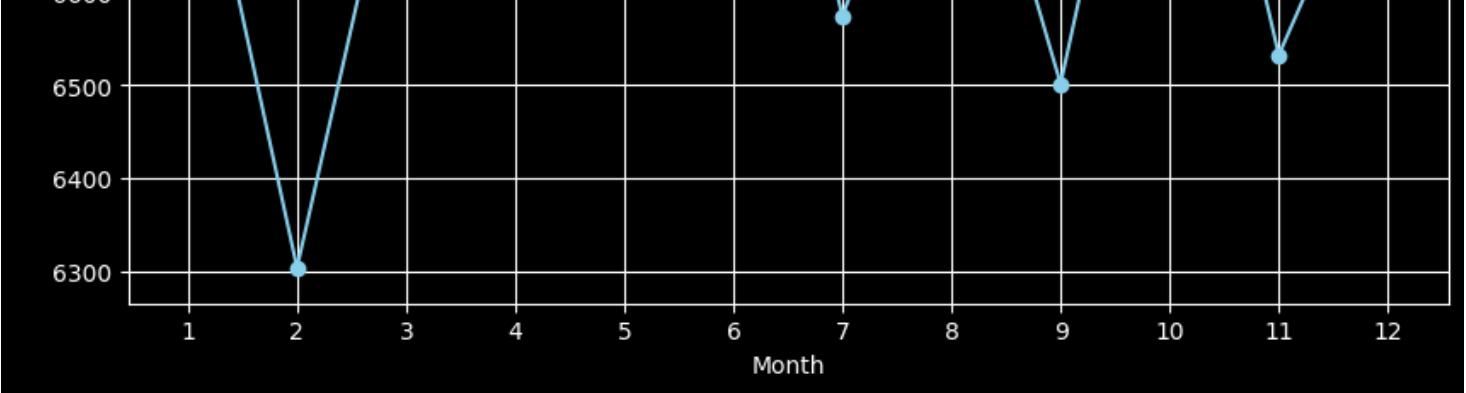


In [318]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Month', 'Total_Spent'], inplace=True)
month_totals = df_copy.groupby('Month')['Total_Spent'].sum()
plt.figure(figsize=(10, 6))
plt.plot(month_totals.index, month_totals.values, marker='o', linestyle='--', color='skyblue')
plt.title('Total Spent by Month')
plt.xlabel('Month')
plt.ylabel('Total Spent')
plt.xticks(month_totals.index)
plt.grid(True)
plt.show()
```

Total Spent by Month





The total spent is not very affected by which month we are in. But it is still considerable.

In [319]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9540 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Transaction_ID  9540 non-null    object 
 1   Item              8613 non-null    object 
 2   Quantity          9086 non-null    object 
 3   Price_Per_Unit   9034 non-null    object 
 4   Total_Spent       9064 non-null    object 
 5   Payment_Method   6525 non-null    object 
 6   Location          5761 non-null    object 
 7   Transaction_Date 9540 non-null    datetime64[ns]
 8   Day               9540 non-null    object 
 9   Month              9540 non-null    int32  
dtypes: datetime64[ns](1), int32(1), object(8)
memory usage: 782.6+ KB
```

In [320]:

```
df.drop(['Location', 'Payment_Method'], axis=1, inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9540 entries, 0 to 9999
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Transaction_ID  9540 non-null    object 
 1   Item              8613 non-null    object 
 2   Quantity          9086 non-null    object 
 3   Price_Per_Unit   9034 non-null    object 
 4   Total_Spent       9064 non-null    object 
 5   Transaction_Date 9540 non-null    datetime64[ns]
 6   Day               9540 non-null    object 
 7   Month              9540 non-null    int32  
dtypes: datetime64[ns](1), int32(1), object(6)
memory usage: 633.5+ KB
```

The following is a code to check price per each item.

In [321]:

```
if 'Price_Per_Unit' in df.columns and 'Item' in df.columns:
    item_prices = df[['Item', 'Price_Per_Unit']].copy()
    item_prices.drop_duplicates(inplace=True)
    item_prices['Price_Per_Unit'] = pd.to_numeric(item_prices['Price_Per_Unit'], errors='coerce')
    item_prices.dropna(subset=['Price_Per_Unit'], inplace=True)
    print("Item Name and Price Per Unit:")
```

```

for index, row in item_prices.iterrows():
    print(f"Item: {row['Item']}, Price Per Unit: {row['Price_Per_Unit']:.2f}")
else:
    print("The dataframe does not contain both 'Item' and 'Price_Per_Unit' columns.")

```

Item Name and Price Per Unit:

Item: Coffee, Price Per Unit: 2.00
 Item: Cake, Price Per Unit: 3.00
 Item: Cookie, Price Per Unit: 1.00
 Item: Salad, Price Per Unit: 5.00
 Item: Smoothie, Price Per Unit: 4.00
 Item: nan, Price Per Unit: 3.00
 Item: Sandwich, Price Per Unit: 4.00
 Item: nan, Price Per Unit: 1.50
 Item: Juice, Price Per Unit: 3.00
 Item: nan, Price Per Unit: 2.00
 Item: nan, Price Per Unit: 1.00
 Item: Tea, Price Per Unit: 1.50
 Item: nan, Price Per Unit: 5.00
 Item: nan, Price Per Unit: 4.00

Price Per Unit of Cookie is uniquely shared with one nan. Therefore we can use this info to fix the nan value that has Price Per Unit = 1. Similarly, we apply the same procedure for Tea and Salad.

In [322]:

```
df.loc[(df['Price_Per_Unit'] == 1) & (df['Item'].isna()), 'Item'] = 'Cookie'
```

In [323]:

```
df.loc[(df['Price_Per_Unit'] == 1.5) & (df['Item'].isna()), 'Item'] = 'Tea'
```

In [324]:

```
df.loc[(df['Price_Per_Unit'] == 5) & (df['Item'].isna()), 'Item'] = 'Salad'
```

In [325]:

```

if 'Price_Per_Unit' in df.columns and 'Item' in df.columns:
    item_prices = df[['Item', 'Price_Per_Unit']].copy()
    item_prices['Price_Per_Unit'] = pd.to_numeric(item_prices['Price_Per_Unit'], errors='coerce')
    item_prices.dropna(subset=['Price_Per_Unit'], inplace=True)
    nan_item_prices = item_prices[item_prices['Item'].isna()]['Price_Per_Unit'].unique()
    for price in nan_item_prices:
        possible_items = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].unique()
        if len(possible_items) > 0:
            item_distribution = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].value_counts(normalize=True)
            nan_indices = df.index[(df['Item'].isna()) & (pd.to_numeric(df['Price_Per_Unit'], errors='coerce') == price)]
            if not item_distribution.empty and len(nan_indices) > 0:
                df.loc[nan_indices, 'Item'] = np.random.choice(item_distribution.index,
                                                               size=len(nan_indices),
                                                               p=item_distribution.values)
        else:
            print(f"Warning: No valid items found for Price_Per_Unit: {price}. NaNs for this price will remain.")
    print("Unique Item Name and Price Per Unit after filling NaNs:")
    item_prices_filled = df[['Item', 'Price_Per_Unit']].copy()
    item_prices_filled['Price_Per_Unit'] = pd.to_numeric(item_prices_filled['Price_Per_Unit'], errors='coerce')
    item_prices_filled.dropna(subset=['Item', 'Price_Per_Unit'], inplace=True)
    item_prices_filled['Item_lower'] = item_prices_filled['Item'].str.lower()
    unique_item_prices = item_prices_filled.drop_duplicates(subset=['Item_lower', 'Price_Per_Unit'])
    for index, row in unique_item_prices.iterrows():
        print(f"Item: {row['Item']}, Price Per Unit: {row['Price_Per_Unit']:.2f}")
else:

```

```
print("The dataframe does not contain both 'Item' and 'Price_Per_Unit' columns.")
```

Unique Item Name and Price Per Unit after filling NaNs:

Item: Coffee, Price Per Unit: 2.00
Item: Cake, Price Per Unit: 3.00
Item: Cookie, Price Per Unit: 1.00
Item: Salad, Price Per Unit: 5.00
Item: Smoothie, Price Per Unit: 4.00
Item: Juice, Price Per Unit: 3.00
Item: Sandwich, Price Per Unit: 4.00
Item: Tea, Price Per Unit: 1.50

Now the items are filled successfully when one of price per unit or item is nan.

In the following we will deal with price per unit and item being both nan at the same time in an instance.

In [326]:

```
df[df['Item'].isna()]
```

Out [326]:

| | Transaction_ID | Item | Quantity | Price_Per_Unit | Total_Spent | Transaction_Date | Day | Month |
|------|----------------|------|----------|----------------|-------------|------------------|-----------|-------|
| 118 | TXN_4633784 | NaN | 5 | NaN | 15.0 | 2023-02-06 | Monday | 2 |
| 151 | TXN_4031509 | NaN | 4 | NaN | 16.0 | 2023-01-04 | Wednesday | 1 |
| 289 | TXN_3495950 | NaN | 4 | NaN | 6.0 | 2023-02-19 | Sunday | 2 |
| 334 | TXN_2523298 | NaN | 4 | NaN | 6.0 | 2023-03-25 | Saturday | 3 |
| 550 | TXN_4186681 | NaN | 4 | NaN | 6.0 | 2023-05-24 | Wednesday | 5 |
| 750 | TXN_5787508 | NaN | 3 | NaN | 9.0 | 2023-07-23 | Sunday | 7 |
| 818 | TXN_7940202 | NaN | 1 | NaN | 4.0 | 2023-07-23 | Sunday | 7 |
| 1154 | TXN_2473090 | NaN | 2 | NaN | 3.0 | 2023-03-03 | Friday | 3 |
| 1337 | TXN_5031214 | NaN | 5 | NaN | 5.0 | 2023-07-29 | Saturday | 7 |
| 1377 | TXN_8396271 | NaN | 2 | NaN | 2.0 | 2023-09-12 | Tuesday | 9 |
| 1589 | TXN_5245399 | NaN | 5 | NaN | 10.0 | 2023-12-25 | Monday | 12 |
| 1761 | TXN_3611851 | NaN | 4 | NaN | NaN | 2023-02-09 | Thursday | 2 |
| 1786 | TXN_1923349 | NaN | 4 | NaN | 6.0 | 2023-07-06 | Thursday | 7 |
| 2002 | TXN_5206049 | NaN | 3 | NaN | 3.0 | 2023-06-24 | Saturday | 6 |
| 2227 | TXN_3200203 | NaN | 2 | NaN | 8.0 | 2023-12-04 | Monday | 12 |
| 2289 | TXN_7524977 | NaN | 4 | NaN | NaN | 2023-12-09 | Saturday | 12 |
| 2596 | TXN_4844386 | NaN | 5 | NaN | 15.0 | 2023-10-28 | Saturday | 10 |
| 2610 | TXN_8266689 | NaN | 5 | NaN | 15.0 | 2023-12-03 | Sunday | 12 |
| 2962 | TXN_9702662 | NaN | 4 | NaN | 16.0 | 2023-07-10 | Monday | 7 |
| 3013 | TXN_1842697 | NaN | 5 | NaN | 15.0 | 2023-10-25 | Wednesday | 10 |
| 3404 | TXN_7797231 | NaN | 5 | NaN | 10.0 | 2023-05-05 | Friday | 5 |
| 3434 | TXN_6457997 | NaN | 1 | NaN | 4.0 | 2023-12-12 | Tuesday | 12 |
| 3666 | TXN_8616276 | NaN | 2 | NaN | 3.0 | 2023-07-22 | Saturday | 7 |
| 3739 | TXN_4849180 | NaN | 5 | NaN | 15.0 | 2023-10-14 | Saturday | 10 |
| 3779 | TXN_7376255 | NaN | NaN | NaN | 25.0 | 2023-05-27 | Saturday | 5 |
| 3900 | TXN_5093855 | NaN | 4 | NaN | 6.0 | 2023-11-23 | Thursday | 11 |
| 4092 | TXN_1840897 | NaN | 1 | NaN | 5.0 | 2023-06-03 | Saturday | 6 |
| 4152 | TXN_9646000 | NaN | 2 | NaN | NaN | 2023-12-14 | Thursday | 12 |
| 4621 | TXN_7844352 | NaN | 2 | NaN | 6.0 | 2023-07-26 | Wednesday | 7 |
| 5039 | TXN_9514452 | NaN | 5 | NaN | 10.0 | 2023-07-29 | Saturday | 7 |

| 5891 | TransactionID | Item | Quantity | Price_Per_Unit | Total_Spent | TransactionDate | Wednesday | Month |
|------|---------------|------|----------|----------------|-------------|-----------------|-----------|-------|
| 5991 | TXN_2913107 | NaN | 4 | NaN | 8.0 | 2023-05-20 | Saturday | 5 |
| 6177 | TXN_3232279 | NaN | 4 | NaN | 16.0 | 2023-05-30 | Tuesday | 5 |
| 6345 | TXN_4208919 | NaN | 3 | NaN | 12.0 | 2023-05-30 | Tuesday | 5 |
| 6429 | TXN_2536573 | NaN | 2 | NaN | 8.0 | 2023-06-24 | Saturday | 6 |
| 6849 | TXN_7928378 | NaN | 3 | NaN | 12.0 | 2023-07-04 | Tuesday | 7 |
| 7478 | TXN_9710982 | NaN | 4 | NaN | 16.0 | 2023-03-03 | Friday | 3 |
| 7597 | TXN_1082717 | NaN | NaN | NaN | 9.0 | 2023-12-13 | Wednesday | 12 |
| 8129 | TXN_8384699 | NaN | 3 | NaN | 15.0 | 2023-12-16 | Saturday | 12 |
| 8458 | TXN_7595907 | NaN | 5 | NaN | 15.0 | 2023-04-24 | Monday | 4 |
| 8463 | TXN_2150872 | NaN | 5 | NaN | 15.0 | 2023-03-18 | Saturday | 3 |
| 8817 | TXN_8365530 | NaN | 4 | NaN | 16.0 | 2023-03-10 | Friday | 3 |
| 8959 | TXN_3803063 | NaN | 4 | NaN | 12.0 | 2023-11-23 | Thursday | 11 |
| 9174 | TXN_5935353 | NaN | 2 | NaN | 8.0 | 2023-05-01 | Monday | 5 |
| 9234 | TXN_8918253 | NaN | 3 | NaN | 12.0 | 2023-09-07 | Thursday | 9 |
| 9425 | TXN_2065683 | NaN | 3 | NaN | 6.0 | 2023-09-21 | Thursday | 9 |
| 9673 | TXN_2480808 | NaN | 1 | NaN | 4.0 | 2023-03-30 | Thursday | 3 |
| 9717 | TXN_3334632 | NaN | 1 | NaN | 2.0 | 2023-11-20 | Monday | 11 |
| 9819 | TXN_1208561 | NaN | NaN | NaN | 20.0 | 2023-08-19 | Saturday | 8 |
| 9820 | TXN_8751702 | NaN | 5 | NaN | 15.0 | 2023-02-13 | Monday | 2 |
| 9996 | TXN_9659401 | NaN | 3 | NaN | 3.0 | 2023-06-02 | Friday | 6 |

In the original notebook, missing features (`Price_Per_Unit`, `Quantity`) were imputed using the target variable (`Total_Spent`). This caused target leakage and an unrealistic R² score of 0.999.

The block below contains the corrected, non-leaky imputation:

1. Convert `Price_Per_Unit`, `Quantity`, and `Total_Spent` to numeric types.
2. Drop any rows where the target (`Total_Spent`) is missing.
3. Impute missing `Price_Per_Unit` and `Quantity` using the median value for that specific `Item`, which is a valid, non-leaky strategy.
4. Handle any remaining NaNs by filling with the global median of the column.

In [327]:

```
# 1. Convert columns to numeric type to be able to use .median()
df['Price_Per_Unit'] = pd.to_numeric(df['Price_Per_Unit'], errors='coerce')
df['Quantity'] = pd.to_numeric(df['Quantity'], errors='coerce')
df['Total_Spent'] = pd.to_numeric(df['Total_Spent'], errors='coerce')

# 2. Handle the Target Variable ('Total_Spent')
# Now we can safely drop rows where the *numeric* target is unknown.
df.dropna(subset=['Total_Spent'], inplace=True)
print(f"Shape after dropping rows with missing 'Total_Spent': {df.shape}")

# 3. Impute 'Price_Per_Unit' Feature (Non-Leaky)
print(f"Missing 'Price_Per_Unit' values before imputation: {df['Price_Per_Unit'].isnull().sum()}")
median_prices = df.groupby('Item')['Price_Per_Unit'].median()
df['Price_Per_Unit'] = df.apply(
    lambda row: median_prices[row['Item']] if pd.isnull(row['Price_Per_Unit']) and pd.notnull(row['Item']) else row['Price_Per_Unit'],
    axis=1
)
# Handle any remaining NaNs (e.g., if an Item had no median price OR the Item was nan)
df['Price_Per_Unit'].fillna(df['Price_Per_Unit'].median(), inplace=True)
print(f"Missing 'Price_Per_Unit' values after imputation: {df['Price_Per_Unit'].isnull().sum()}"
```

```
m() })
```

```
# 4. Impute 'Quantity' Feature (Non-Leaky)
print(f"Missing 'Quantity' values before imputation: {df['Quantity'].isnull().sum()}")
median_quantities = df.groupby('Item')['Quantity'].median()
df['Quantity'] = df.apply(
    lambda row: median_quantities[row['Item']] if pd.isnull(row['Quantity']) and pd.notnull(row['Item']) else row['Quantity'],
    axis=1
)
# Handle any remaining NaNs (e.g., if an Item had no median price OR the Item was nan)
df['Quantity'].fillna(df['Quantity'].median(), inplace=True)
print(f"Missing 'Quantity' values after imputation: {df['Quantity'].isnull().sum()}")
```

Shape after dropping rows with missing 'Total_Spent': (9064, 8)

Missing 'Price_Per_Unit' values before imputation: 487

Missing 'Price_Per_Unit' values after imputation: 0

Missing 'Quantity' values before imputation: 434

Missing 'Quantity' values after imputation: 0

```
/tmp/ipython-input-3580754496.py:19: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Price_Per_Unit'].fillna(df['Price_Per_Unit'].median(), inplace=True)
```

```
/tmp/ipython-input-3580754496.py:30: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Quantity'].fillna(df['Quantity'].median(), inplace=True)
```

All remaining instances have 2 NaN that can't be replaced by an exact value 100%. Therefore, it is better if we remove it but first we will try to replace NaN in item.

In [328]:

```
df.loc[(df['Price_Per_Unit'] == 5) & (df['Item'].isna()), 'Item'] = 'Salad'
df.loc[(df['Price_Per_Unit'] == 1.5) & (df['Item'].isna()), 'Item'] = 'Tea'
df.loc[(df['Price_Per_Unit'] == 1) & (df['Item'].isna()), 'Item'] = 'Cookie'
if 'Price_Per_Unit' in df.columns and 'Item' in df.columns:
    item_prices = df[['Item', 'Price_Per_Unit']].copy()
    item_prices['Price_Per_Unit'] = pd.to_numeric(item_prices['Price_Per_Unit'], errors='coerce')
    item_prices.dropna(subset=['Price_Per_Unit'], inplace=True)
    nan_item_prices = item_prices[item_prices['Item'].isna()]['Price_Per_Unit'].unique()
    for price in nan_item_prices:
        possible_items = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].unique()
        if len(possible_items) > 0:
            item_distribution = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].value_counts(normalize=True)
            nan_indices = df.index[(df['Item'].isna()) & (pd.to_numeric(df['Price_Per_Unit'], errors='coerce') == price)]
            if not item_distribution.empty and len(nan_indices) > 0:
                df.loc[nan_indices, 'Item'] = np.random.choice(item_distribution.index,
                                                               size=len(nan_indices),
                                                               p=item_distribution.values)
)
else:
    print(f"Warning: No valid items found for Price_Per_Unit: {price}. NaNs for t
```

```

his price will remain.")

print("Unique Item Name and Price Per Unit after filling NaNs:")
item_prices_filled = df[['Item', 'Price_Per_Unit']].copy()
item_prices_filled['Price_Per_Unit'] = pd.to_numeric(item_prices_filled['Price_Per_Unit'], errors='coerce')
item_prices_filled.dropna(subset=['Item', 'Price_Per_Unit'], inplace=True)
item_prices_filled['Item_lower'] = item_prices_filled['Item'].str.lower()
unique_item_prices = item_prices_filled.drop_duplicates(subset=['Item_lower', 'Price_Per_Unit'])

for index, row in unique_item_prices.iterrows():
    print(f"Item: {row['Item']}, Price Per Unit: {row['Price_Per_Unit']:.2f}")

else:
    print("The dataframe does not contain both 'Item' and 'Price_Per_Unit' columns.")

```

Unique Item Name and Price Per Unit after filling NaNs:

Item: Coffee, Price Per Unit: 2.00
 Item: Cake, Price Per Unit: 3.00
 Item: Salad, Price Per Unit: 5.00
 Item: Smoothie, Price Per Unit: 4.00
 Item: Juice, Price Per Unit: 3.00
 Item: Sandwich, Price Per Unit: 4.00
 Item: Cookie, Price Per Unit: 1.00
 Item: Tea, Price Per Unit: 1.50

In [329]:

```

df.dropna(inplace=True)
print("DataFrame Info:")
df.info()
print("\nNaN values per column:")
print(df.isnull().sum())

```

DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
Index: 9064 entries, 0 to 9999
Data columns (total 8 columns):
Column Non-Null Count Dtype

0 Transaction_ID 9064 non-null object
1 Item 9064 non-null object
2 Quantity 9064 non-null float64
3 Price_Per_Unit 9064 non-null float64
4 Total_Spent 9064 non-null float64
5 Transaction_Date 9064 non-null datetime64[ns]
6 Day 9064 non-null object
7 Month 9064 non-null int32
dtypes: datetime64[ns](1), float64(3), int32(1), object(3)
memory usage: 859.9+ KB

Nan values per column:
Transaction_ID 0
Item 0
Quantity 0
Price_Per_Unit 0
Total_Spent 0
Transaction_Date 0
Day 0
Month 0
dtype: int64

Now data is fully preprocessed and 100% clean, with no leakage. I can now run the "EDA after data preprocessing" section.

EDA after data preprocessing

In [330]:

```

print("\nDataFrame Description:")

```

```
print("DataFrame Description: ")
print(df.describe(include='all'))
```

DataFrame Description:

| | Transaction_ID | Item | Quantity | Price_Per_Unit | Total_Spent | \ |
|--------|----------------|-------|-------------|----------------|-------------|---|
| count | 9064 | 9064 | 9064.000000 | 9064.000000 | 9064.000000 | |
| unique | 9064 | 8 | NaN | NaN | NaN | |
| top | TXN_6170729 | Juice | NaN | NaN | NaN | |
| freq | 1 | 1203 | NaN | NaN | NaN | |
| mean | NaN | NaN | 3.022727 | 2.946988 | 8.914442 | |
| min | NaN | NaN | 1.000000 | 1.000000 | 1.000000 | |
| 25% | NaN | NaN | 2.000000 | 2.000000 | 4.000000 | |
| 50% | NaN | NaN | 3.000000 | 3.000000 | 8.000000 | |
| 75% | NaN | NaN | 4.000000 | 4.000000 | 12.000000 | |
| max | NaN | NaN | 5.000000 | 5.000000 | 25.000000 | |
| std | NaN | NaN | 1.384699 | 1.276373 | 6.010311 | |

| | Transaction_Date | Day | Month |
|--------|-------------------------------|----------|-------------|
| count | 9064 | 9064 | 9064.000000 |
| unique | NaN | 7 | NaN |
| top | NaN | Thursday | NaN |
| freq | NaN | 1319 | NaN |
| mean | 2023-07-02 06:26:03.283318272 | NaN | 6.534753 |
| min | 2023-01-01 00:00:00 | NaN | 1.000000 |
| 25% | 2023-04-01 00:00:00 | NaN | 4.000000 |
| 50% | 2023-07-02 00:00:00 | NaN | 7.000000 |
| 75% | 2023-10-02 00:00:00 | NaN | 10.000000 |
| max | 2023-12-31 00:00:00 | NaN | 12.000000 |
| std | NaN | NaN | 3.445499 |

In [331]:

```
item_total_spent = df.groupby('Item')['Total_Spent'].sum().sort_values(ascending=False)
plt.figure(figsize=(12, 6))
sns.barplot(x=item_total_spent.index, y=item_total_spent.values)
plt.xticks(rotation=90)
plt.title('Total Spent by Item')
plt.xlabel('Item')
plt.ylabel('Total Spent')
plt.tight_layout()
plt.show()

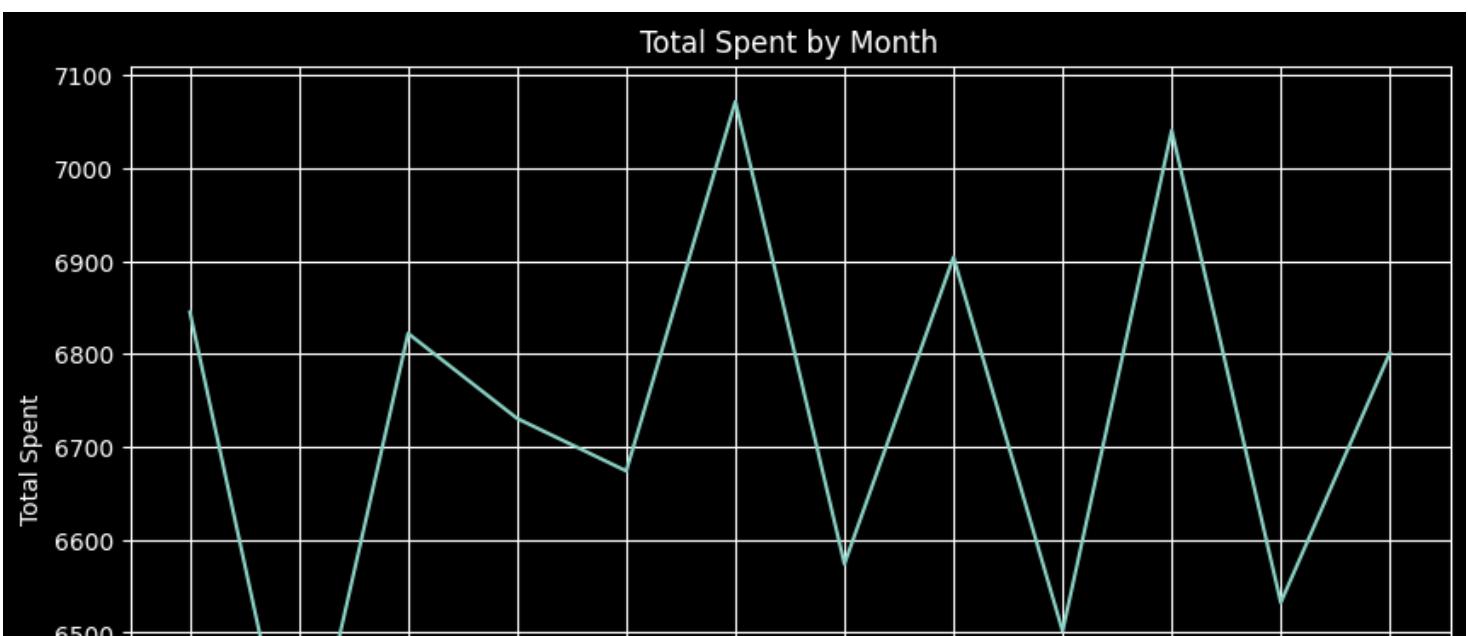
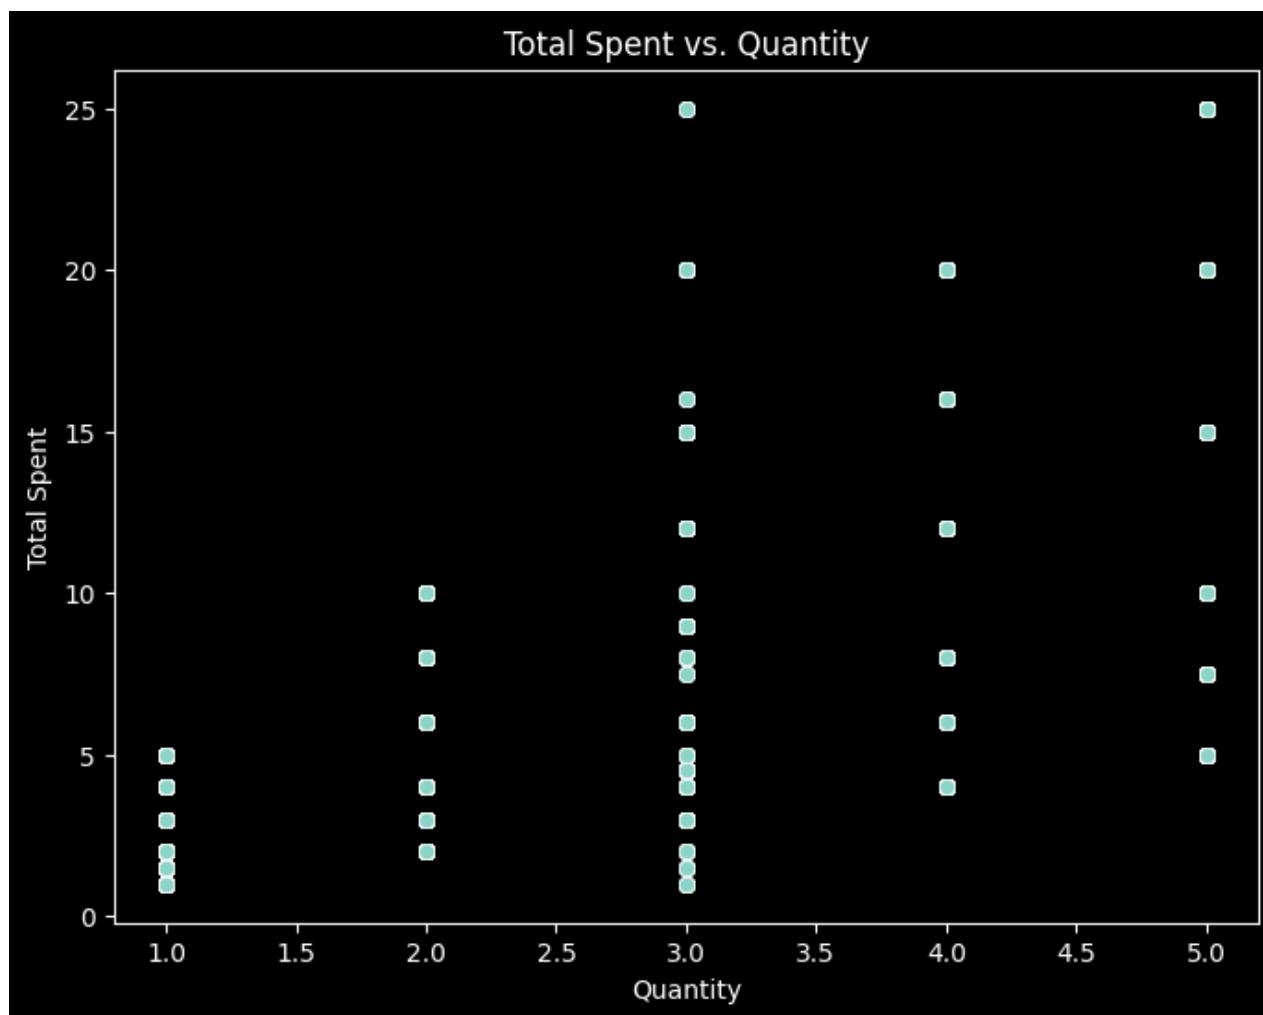
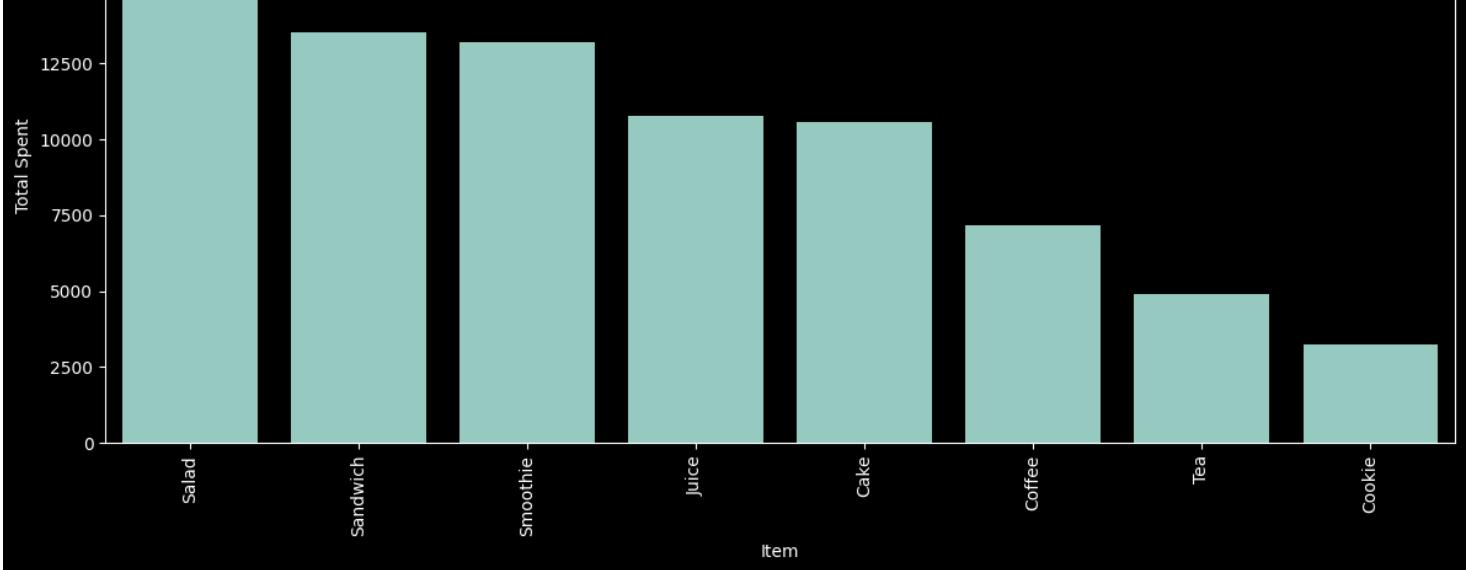
plt.figure(figsize=(8, 6))
sns.scatterplot(x=df['Quantity'], y=df['Total_Spent'])
plt.title('Total Spent vs. Quantity')
plt.xlabel('Quantity')
plt.ylabel('Total Spent')
plt.show()

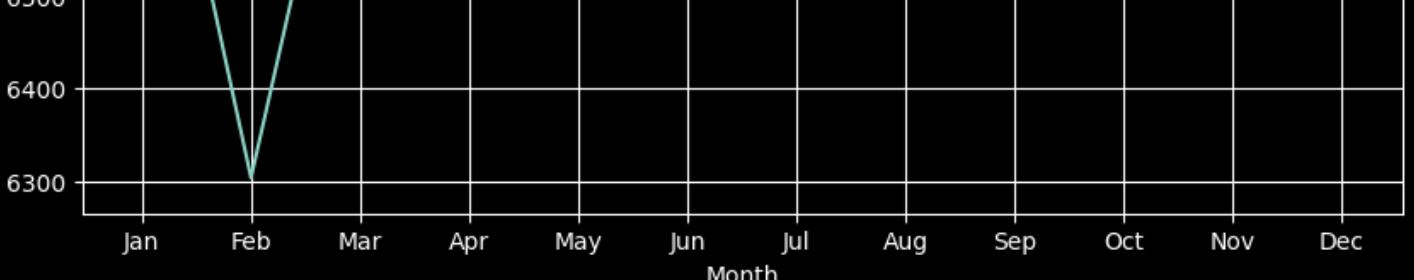
month_total_spent = df.groupby('Month')['Total_Spent'].sum()
plt.figure(figsize=(10, 6))
sns.lineplot(x=month_total_spent.index, y=month_total_spent.values)
plt.title('Total Spent by Month')
plt.xlabel('Month')
plt.ylabel('Total Spent')
plt.xticks(month_total_spent.index, ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']) # Label months
plt.grid(True)
plt.show()

df['Day_Name'] = df['Transaction_Date'].dt.day_name()
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
day_name_total_spent = df.groupby('Day_Name')['Total_Spent'].sum().reindex(day_order)
plt.figure(figsize=(10, 6))
sns.lineplot(x=day_name_total_spent.index, y=day_name_total_spent.values)
plt.title('Total Spent by Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Total Spent')
plt.grid(True)
plt.tight_layout()
plt.show()
```

Total Spent by Item







Total Spent by Day of Week

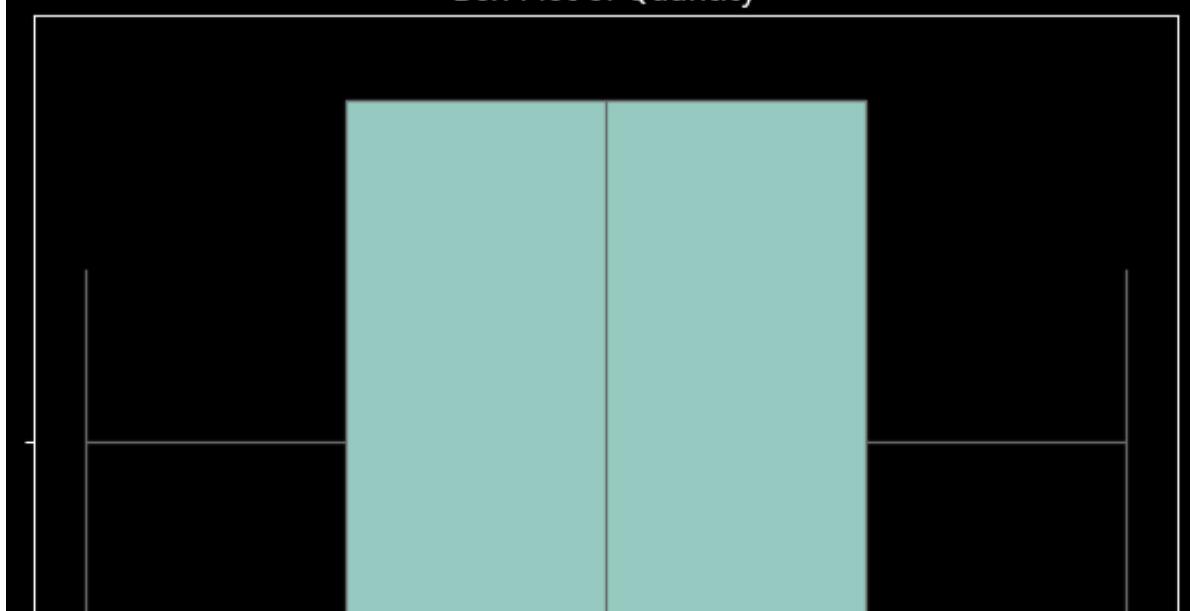


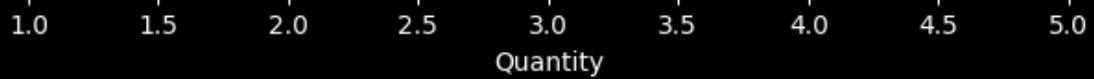
In [332]:

```
numerical_cols = ['Quantity', 'Price_Per_Unit', 'Day_Name', 'Month']

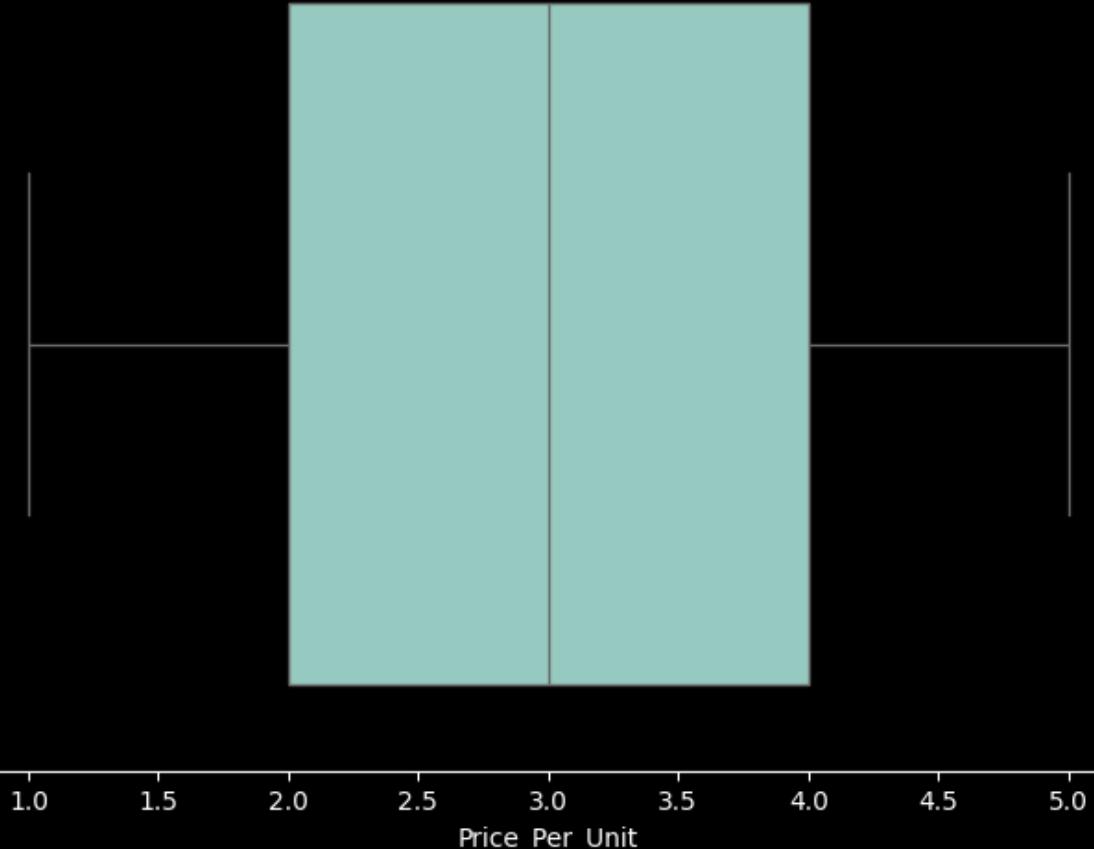
for col in numerical_cols:
    if col in df.columns:
        plt.figure(figsize=(8, 6))
        sns.boxplot(x=df[col])
        plt.title(f'Box Plot of {col}')
        plt.xlabel(col)
        plt.show()
```

Box Plot of Quantity

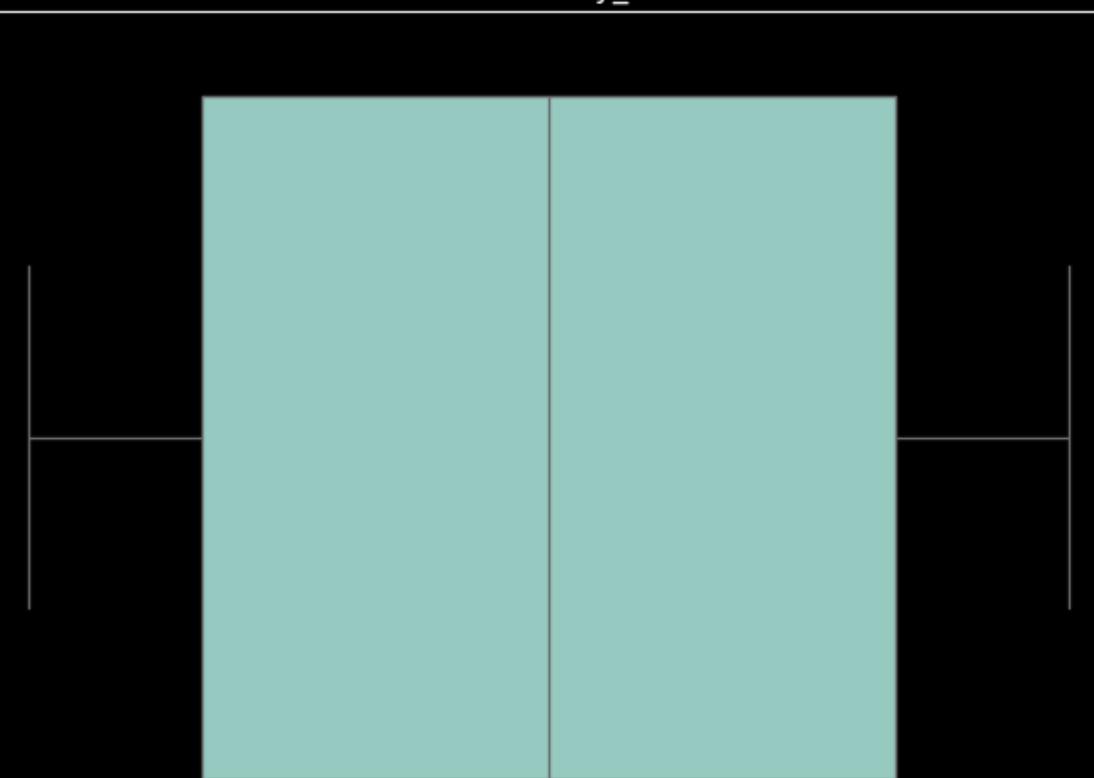




Box Plot of Price_Per_Unit

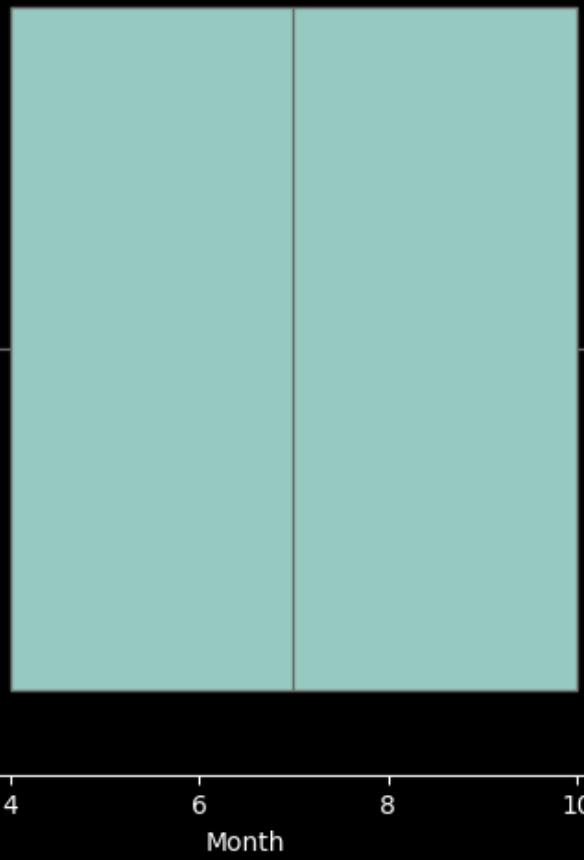


Box Plot of Day_Name



| | | | | | | |
|----------|---------|----------|--------|----------|-----------|--------|
| Friday | Tuesday | Thursday | Sunday | Saturday | Wednesday | Monday |
| Day_Name | | | | | | |

Box Plot of Month



Dataset Balance

Since the dataset lacks timestamps (no hours), we cannot analyze hourly trends (so “peak hours” analysis is skipped). Therefore, we will plot our target (total spent) that will be predicted later on.

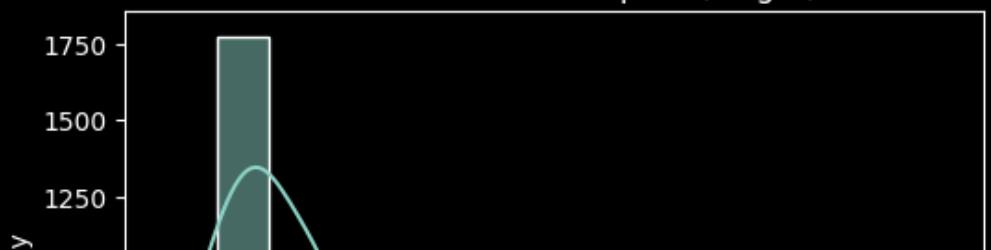
-

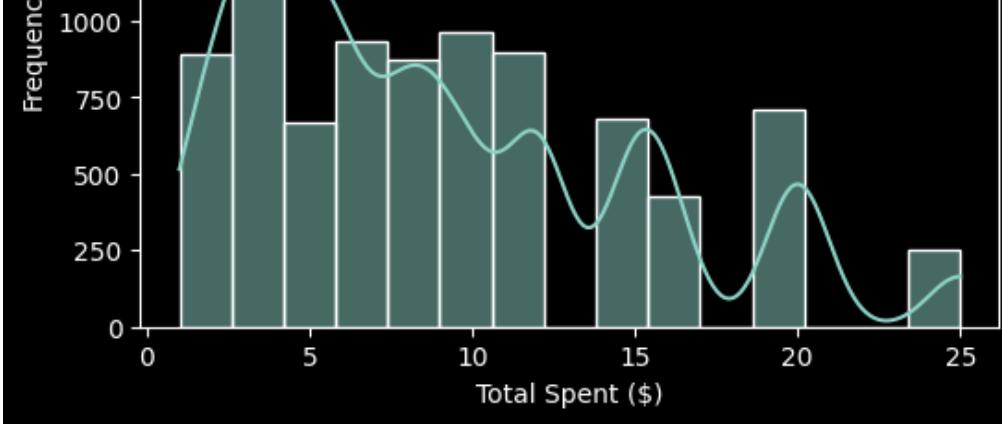
Since this is regression, we inspect *balance* by plotting a histogram (or KDE) of Total_Spent. A highly skewed target (heavy outliers) would suggest careful metric choice.

In [333] :

```
y = df['Total_Spent'] # target variable
# histogram of Total_Spent
plt.figure(figsize=(6,4))
sns.histplot(y, bins=15, kde=True)
plt.title('Distribution of Total Spent (Target)')
plt.xlabel('Total Spent ($)')
plt.ylabel('Frequency')
plt.show()
```

Distribution of Total Spent (Target)





The range is between 1 and 25.

The median is around 7.5 and mean about 8.8.

It is clear that we are dealing with *positive-to-zero* skewness --> our target is nicely distributed around a central value. Therefore standard regression metrics must be used (MAE,MSE,RMSE,R²).

Regression Metrics

1. Because this is a regression problem (continuous target), we will focus on regression metrics like MAE, MSE, RMSE, and R².
2. Confusion matrix is not applicable here, as confusion matrix summarizes classification errors, not continuous regression predictions.
3. RMSE is in the same units as sales (\$) and penalizes large errors more than MAE does.
4. R² of 1.0 means perfect fit, whereas 0 means no predictive power beyond the mean.
5. Quantile (pinball) loss is not about where most of your data live—it's about which “slice” of the error distribution you care to optimize. In my case, I decided to optimize big spenders, so I choose $\alpha=0.9$.

Standard Train-Test Split

Setting a random seed ensures that anyone running this code will get the same split and results.

In [334]:

```
from sklearn.model_selection import train_test_split

# these columns could be easily removed because I already extracted Day and Month columns
df.drop(columns=['Transaction_Date', 'Transaction_ID'], inplace=True)

# target variable
y = df['Total_Spent']

# categorical variables (Item, Day) are one-hot encoded because computers understand numbers
# better than string (it will be used in Feature selection)
# used one-hot encoding to prevent association.
X = pd.get_dummies(df.drop(columns=['Total_Spent']), columns=['Item', 'Day'], drop_first=True)
X.drop(columns=['Day_Name'], inplace=True)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.4, random_state=42
)

# printing the shapes is used to confirm the size of each
print("Train set:", X_train.shape, "Test set:", X_test.shape)
```

Train set: (5438, 16) Test set: (3626, 16)

In [335]:

```
X.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 9064 entries, 0 to 9999
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Quantity         9064 non-null    float64
 1   Price_Per_Unit  9064 non-null    float64
 2   Month            9064 non-null    int32  
 3   Item_Coffee      9064 non-null    bool   
 4   Item_Cookie      9064 non-null    bool   
 5   Item_Juice       9064 non-null    bool   
 6   Item_Salad       9064 non-null    bool   
 7   Item_Sandwich    9064 non-null    bool   
 8   Item_Smoothie   9064 non-null    bool   
 9   Item_Tea          9064 non-null    bool   
 10  Day_Monday      9064 non-null    bool   
 11  Day_Saturday    9064 non-null    bool   
 12  Day_Sunday       9064 non-null    bool   
 13  Day_Thursday     9064 non-null    bool   
 14  Day_Tuesday      9064 non-null    bool   
 15  Day_Wednesday   9064 non-null    bool   
dtypes: bool(13), float64(2), int32(1)
memory usage: 621.0 KB

```

Feature Selection

In the following we perform feature selection to reduce dimensionality:

1. SelectKBest picks the top k features by a statistical test.
2. RFE (Recursive Feature Elimination) uses a model feature importances to iteratively drop the least important features.
3. Sequential Forward/Backward Selection would add/remove features one by one.

I will not try Sequential Selection because it can be very slow.

In [336]:

```

from sklearn.feature_selection import SelectKBest, f_regression
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

# use SelectKBest with f_regression to select top 5 features
skb = SelectKBest(score_func=f_regression, k=5)
X_train_skb = skb.fit_transform(X_train, y_train)
X_test_skb = skb.transform(X_test)
print("After SelectKBest, Train shape:", X_train_skb.shape, "      Test shape:", X_test_skb.shape)

# use RFE with a linear model to select top 5 features
rfe = RFE(estimator=LinearRegression(), n_features_to_select=5)
X_train_rfe = rfe.fit_transform(X_train, y_train)
X_test_rfe = rfe.transform(X_test)
print("After RFE, Train shape:", X_train_rfe.shape, "      Test shape:", X_test_rfe.shape)

```

```

After SelectKBest, Train shape: (5438, 5)      Test shape: (3626, 5)
After RFE, Train shape: (5438, 5)      Test shape: (3626, 5)

```

Model Building and Evaluation/Comparison

Each model is trained on the same train set and evaluated on the same test set. We record MAE, MSE, RMSE, and R² for comparison.

In [337]:

```

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, mean_pinball_loss

```

```

from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import GradientBoostingRegressor

models = {
    "Linear Regression": LinearRegression(),
    "MLP Regressor": MLPRegressor(hidden_layer_sizes=(200,), max_iter=300, random_state=42),
    "KNN Regressor": KNeighborsRegressor(n_neighbors=5),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

alpha = 0.9 # 90th percentile pinball loss

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
    r2 = r2_score(y_test, preds)
    qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

    results.append([name, mae, mse, rmse, r2, qloss])

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))

```

| | Model | MAE | MSE | RMSE | R2 | \ |
|---|--------------------|----------|----------|----------|----------|---|
| 0 | Linear Regression | 1.440077 | 4.246647 | 2.060739 | 0.882051 | |
| 1 | MLP Regressor | 0.332508 | 1.089320 | 1.043705 | 0.969744 | |
| 2 | KNN Regressor | 0.710618 | 1.821051 | 1.349463 | 0.949421 | |
| 3 | Gradient Boosting | 0.263807 | 1.050543 | 1.024960 | 0.970821 | |
| | PinballLoss(q=0.9) | | | | | |
| 0 | | 0.692072 | | | | |
| 1 | | 0.128453 | | | | |
| 2 | | 0.334316 | | | | |
| 3 | | 0.119858 | | | | |

Gradient Boosting (R²: 0.97) and MLP Regressor (R²: 0.97) are the top performers. This is expected, as tree-based models and neural networks are very effective at capturing the non-linear interactions between the two most important features: *Quantity* and *Price_Per_Unit*.

Linear Regression (R²: 0.88) performs the worst. This is because it can only find a linear relationship, but the true relationship (*Total_Spent* = *Quantity* *Price_Per_Unit*) is multiplicative (non-linear)*.

KNN Regressor (R²: 0.95) does very well by finding "similar" sales in the training data. Its performance is high but dependent on the choice of k (neighbors).

Hyperparameter tuning

In [338]:

```

from sklearn.model_selection import GridSearchCV
"""

from sklearn.linear_model import Ridge
ridge = Ridge()
ridge_params = {'alpha': [0.01, 0.1, 1, 10, 100, 1000]}
ridge_grid = GridSearchCV(ridge, ridge_params, cv=5, scoring='r2')
ridge_grid.fit(X_train, y_train)
print("Best Ridge Params:", ridge_grid.best_params_)
print("")
"""

from sklearn.linear_model import Lasso

```

```

lasso = Lasso(max_iter=10000)
lasso_params = {'alpha': [0.001, 0.01, 0.1, 1, 10, 100, 1000]}
lasso_grid = GridSearchCV(lasso, lasso_params, cv=5, scoring='r2')
lasso_grid.fit(X_train, y_train)
print("Best Lasso Params:", lasso_grid.best_params_)

"""
from sklearn.neural_network import MLPRegressor
mlp = MLPRegressor(max_iter=1000, random_state=42)
mlp_params = {
    'hidden_layer_sizes': [(50,), (100,), (50,50)],
    'activation': ['relu', 'tanh'],
    'solver': ['adam', 'sgd'],
    'alpha': [0.0001, 0.001], # regularization
}
mlp_grid = GridSearchCV(mlp, mlp_params, cv=5, scoring='r2', n_jobs=-1)
mlp_grid.fit(X_train, y_train)
print("Best MLP Params:", mlp_grid.best_params_)
print("")
"""

from sklearn.neighbors import KNeighborsRegressor
knn = KNeighborsRegressor()
knn_params = {
    'n_neighbors': [3, 5, 7, 10],
    'weights': ['uniform', 'distance'],
    'p': [1, 2] # 1: Manhattan, 2: Euclidean
}
knn_grid = GridSearchCV(knn, knn_params, cv=5, scoring='r2')
knn_grid.fit(X_train, y_train)
print("Best KNN Params:", knn_grid.best_params_)

```

```

Best Lasso Params: {'alpha': 0.01}
Best KNN Params: {'n_neighbors': 10, 'p': 2, 'weights': 'uniform'}

```

In [339]:

```

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, mean_pinball_loss
from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import GradientBoostingRegressor

models = {
    "Linear Regression": LinearRegression(),
    "Lasso Regression": Lasso(alpha=0.01),
    "KNN Regressor": KNeighborsRegressor(n_neighbors=5),
    "KNN RegressorHT": KNeighborsRegressor(n_neighbors=10, weights='distance', p=2),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

alpha = 0.9 # 90th percentile pinball loss

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
    r2 = r2_score(y_test, preds)
    qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

    results.append([name, mae, mse, rmse, r2, qloss])

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))

```

| Model | MAE | MSE | RMSE | R2 | \ |
|-------------------|-----|-----|------|----|---|
| Linear Regression | | | | | |
| Lasso Regression | | | | | |
| KNN Regressor | | | | | |
| KNN RegressorHT | | | | | |
| Gradient Boosting | | | | | |

```
0 Linear Regression 1.440077 4.246647 2.060739 0.882051
1 Lasso Regression 1.427307 4.231398 2.057036 0.882474
2 KNN Regressor 0.710618 1.821051 1.349463 0.949421
3 KNN RegressorHT 0.427399 1.793054 1.339050 0.950198
4 Gradient Boosting 0.263807 1.050543 1.024960 0.970821
```

```
PinballLoss(q=0.9)
0 0.692072
1 0.686265
2 0.334316
3 0.202720
4 0.119858
```

In [340]:

```
from sklearn.linear_model import Lasso
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform, randint

lasso_params = {
    'alpha': uniform(0.0001, 1) # Try small to moderate penalties
}
lasso_search = RandomizedSearchCV(
    Lasso(max_iter=10000), # Ensure enough iterations for convergence
    param_distributions=lasso_params,
    n_iter=20,
    cv=5,
    scoring='r2',
    random_state=42
)
lasso_search.fit(X_train, y_train)
print("Lasso Best Params:", lasso_search.best_params_)

from sklearn.neighbors import KNeighborsRegressor
knn_params = {
    'n_neighbors': randint(3, 20),
    'weights': ['uniform', 'distance'],
    'p': [1, 2]
}
knn_search = RandomizedSearchCV(
    KNeighborsRegressor(),
    knn_params,
    n_iter=20,
    cv=5,
    scoring='r2',
    random_state=42
)
knn_search.fit(X_train, y_train)
print("KNN Best Params:", knn_search.best_params_)
```

Lasso Best Params: {'alpha': np.float64(0.020684494295802446)}

KNN Best Params: {'n_neighbors': 17, 'p': 2, 'weights': 'uniform'}

In [341]:

```
models = {
    "Linear Regression": LinearRegression(),
    "Lasso Regression": Lasso(alpha=0.02),
    "KNN Regressor": KNeighborsRegressor(n_neighbors=5),
    "KNN RegressorHT": KNeighborsRegressor(n_neighbors=16, weights='distance', p=2),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

alpha = 0.9 # 90th percentile pinball loss

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
```

```

r2 = r2_score(y_test, preds)
qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

results.append([name, mae, mse, rmse, r2, qloss])

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))

```

| | Model | MAE | MSE | RMSE | R2 | \ |
|---|--------------------|----------|----------|----------|----------|---|
| 0 | Linear Regression | 1.440077 | 4.246647 | 2.060739 | 0.882051 | |
| 1 | Lasso Regression | 1.427166 | 4.229183 | 2.056498 | 0.882536 | |
| 2 | KNN Regressor | 0.710618 | 1.821051 | 1.349463 | 0.949421 | |
| 3 | KNN RegressorHT | 0.417100 | 1.764332 | 1.328281 | 0.950996 | |
| 4 | Gradient Boosting | 0.263807 | 1.050543 | 1.024960 | 0.970821 | |
| | PinballLoss(q=0.9) | | | | | |
| 0 | | 0.692072 | | | | |
| 1 | | 0.686463 | | | | |
| 2 | | 0.334316 | | | | |
| 3 | | 0.198430 | | | | |
| 4 | | 0.119858 | | | | |

The results showed that Random Search outperformed Grid Search for the KNN Regressor, while Grid Search provided better results for Lasso Regression.

KNN involves multiple hyperparameters and less sensitivity to each one (suitable for Random Search which explores large spaces without testing all possibilities/combinations).

The following is another way to see Hyperparameter Tuning if you did not understand the previous results

In [342]:

```

# We will now tune our models to see if we can improve performance.
# We will use GridSearchCV for a smaller search space and RandomizedSearchCV
# for a larger one.

# --- 1. GridSearchCV ---
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Lasso
from sklearn.neighbors import KNeighborsRegressor

print("--- Starting GridSearchCV ---")

# Lasso (Grid)
lasso = Lasso(max_iter=10000)
lasso_params = {'alpha': [0.001, 0.01, 0.1, 1]}
lasso_grid = GridSearchCV(lasso, lasso_params, cv=5, scoring='r2')
lasso_grid.fit(X_train, y_train)
print("Best Lasso Params (Grid):", lasso_grid.best_params_)

# KNN (Grid)
knn = KNeighborsRegressor()
knn_params = {
    'n_neighbors': [3, 5, 7, 10],
    'weights': ['uniform', 'distance'],
    'p': [1, 2] # 1: Manhattan, 2: Euclidean
}
knn_grid = GridSearchCV(knn, knn_params, cv=5, scoring='r2')
knn_grid.fit(X_train, y_train)
print("Best KNN Params (Grid):", knn_grid.best_params_)

# --- 2. RandomizedSearchCV ---
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform, randint

print("\n--- Starting RandomizedSearchCV ---")

# Lasso (Random)
lasso_params_rand = {

```

```

'alpha': uniform(0.0001, 1) # Try small to moderate penalties
}
lasso_search = RandomizedSearchCV(
    Lasso(max_iter=10000),
    param_distributions=lasso_params_rand,
    n_iter=20, cv=5, scoring='r2', random_state=42
)
lasso_search.fit(X_train, y_train)
print("Best Lasso Params (Random):", lasso_search.best_params_)

# KNN (Random)
knn_params_rand = {
    'n_neighbors': randint(3, 20),
    'weights': ['uniform', 'distance'],
    'p': [1, 2]
}
knn_search = RandomizedSearchCV(
    KNeighborsRegressor(),
    knn_params_rand,
    n_iter=20, cv=5, scoring='r2', random_state=42
)
knn_search.fit(X_train, y_train)
print("Best KNN Params (Random):", knn_search.best_params_)
print ("Ended hyperparameter tuning by finding best parameters")
print ("")

```

"""

Final Model Comparison

We will now compare our initial baselines with the best-tuned models from our Grid and Random searches.

Best params found:

```

Lasso (Random): {'alpha': 0.0206...} -> Use 0.02
KNN (Random): {'n_neighbors': 17, 'p': 2, 'weights': 'uniform'}
KNN (Grid): {'n_neighbors': 10, 'p': 2, 'weights': 'uniform'}
Note: Manual test of {'n_neighbors': 16, 'weights': 'distance'} was also very good.
"""

```

```

models = {
    "Linear Regression": LinearRegression(),
    "Lasso (Tuned)": Lasso(alpha=0.02, max_iter=10000), # From RandomSearch
    "KNN (Base)": KNeighborsRegressor(n_neighbors=5),
    "KNN (Tuned)": KNeighborsRegressor(n_neighbors=17, weights='uniform', p=2), # From RandomSearch
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

```

alpha = 0.9 # 90th percentile pinball loss

```

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)

    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
    r2 = r2_score(y_test, preds)
    qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

    results.append([name, mae, mse, rmse, r2, qloss])

```

```

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))

```

--- Starting GridSearchCV ---

Best Lasso Params (Grid): {'alpha': 0.01}

Best KNN Params (Grid): {'n_neighbors': 10, 'p': 2, 'weights': 'uniform'}

--- Starting RandomizedSearchCV ---

Best Lasso Params (Random): {'alpha': np.float64(0.020684494295802446)}

Best KNN Params (Random): {'n_neighbors': 17, 'p': 2, 'weights': 'uniform'}

Ended hyperparameter tuning by finding best parameters

| | Model | MAE | MSE | RMSE | R2 | \ |
|---|-------------------|----------|----------|----------|----------|---|
| 0 | Linear Regression | 1.440077 | 4.246647 | 2.060739 | 0.882051 | |
| 1 | Lasso (Tuned) | 1.427166 | 4.229183 | 2.056498 | 0.882536 | |
| 2 | KNN (Base) | 0.710618 | 1.821051 | 1.349463 | 0.949421 | |
| 3 | KNN (Tuned) | 0.677152 | 1.548526 | 1.244398 | 0.956990 | |
| 4 | Gradient Boosting | 0.263807 | 1.050543 | 1.024960 | 0.970821 | |

PinballLoss (q=0.9)

| | |
|---|----------|
| 0 | 0.692072 |
| 1 | 0.686463 |
| 2 | 0.334316 |
| 3 | 0.323820 |
| 4 | 0.119858 |

Conclusion

We successfully implemented a full end-to-end ML pipeline to predict cafe sales.

Data Processing: The pipeline involved robust data cleaning (to handle 'ERROR' values), feature engineering (extracting Day and Month), and a non-leaky imputation strategy to fix missing data.

Modeling: We compared multiple regressors and found that non-linear models (Gradient Boosting, MLP, KNN) significantly outperformed linear models. This confirms that the relationship between the features (especially Quantity and Price_Per_Unit) and the target (Total_Spent) is not linear.

Tuning: Hyperparameter tuning with RandomizedSearchCV provided a slight improvement over GridSearchCV, especially for the KNN model.

The final Gradient Boosting model was the most accurate, achieving an R² of 0.97 on the test set, demonstrating a very strong predictive model for cafe sales.