

# Cafe Sales Prediction

In [808]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

plt.style.use('dark_background')
df = pd.read_csv('/content/dirty_cafe_sales.csv')
df.head(100)
```

Out[808]:

	Transaction ID	Item	Quantity	Price Per Unit	Total Spent	Payment Method	Location	Transaction Date
0	TXN_1961373	Coffee	2	2.0	4.0	Credit Card	Takeaway	2023-09-08
1	TXN_4977031	Cake	4	3.0	12.0	Cash	In-store	2023-05-16
2	TXN_4271903	Cookie	4	1.0	ERROR	Credit Card	In-store	2023-07-19
3	TXN_7034554	Salad	2	5.0	10.0	UNKNOWN	UNKNOWN	2023-04-27
4	TXN_3160411	Coffee	2	2.0	4.0	Digital Wallet	In-store	2023-06-11
...	...	...	...	...	...	...	...	...
95	TXN_8268061	Salad	3	5.0	15.0	ERROR	Takeaway	2023-08-20
96	TXN_5220895	Salad	5	5.0	25.0	Cash	In-store	2023-06-10
97	TXN_3085509	Coffee	4	2.0	8.0	Digital Wallet	In-store	2023-04-15
98	TXN_9999113	Juice	4	3.0	12.0	Cash	Takeaway	2023-05-27
99	TXN_8779771	Coffee	4	2.0	8.0	Cash	In-store	2023-07-25

100 rows × 8 columns

## Data Preprocessing and EDA

In [809]:

```
import numpy as np
df.replace(['ERROR', 'UNKNOWN'], np.nan, inplace=True)
```


In [810]:

```
df.columns = df.columns.str.replace(' ', '_')
```

In [811]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Payment_Method', 'Total_Spent'], inplace=True)
payment_method_totals = df_copy.groupby('Payment_Method')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(payment_method_totals, labels=payment_method_totals.index,
                                    autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Payment Method')
plt.legend(wedges, payment_method_totals.index, title="Payment Method", loc="center left",
           bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```


Total Spent by Payment Method




After analyzing the distribution of the Payment Method feature, we observe that the total spending is almost evenly divided among the three methods (approximately 33%, 33.5%, and 33.5%). This indicates that the Payment Method does not significantly impact the target variable (Total Spent), and therefore, it can be considered for removal as a non-informative feature.

In [812]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Location', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Location')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Location')
plt.legend(wedges, location_totals.index, title="Location", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```







The analysis of the Location feature shows that total spending is nearly evenly split between In-Store (50.7%) and Takeaway (49.3%). This minimal difference suggests that Location has little to no impact on the target variable (Total Spent), and thus, it can be considered for removal as an uninformative feature.

In [813]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Item', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Item')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Item')
plt.legend(wedges, location_totals.index, title="Item", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```

Total Spent by Item






We observe that there are 8 different items, and the percentage of total spending across these items varies significantly — ranging from 4% to 21.7%. This high variability indicates that the Item feature plays an important role in our analysis and should be retained for further study.

In [814]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Quantity', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Quantity')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Quantity')
plt.legend(wedges, location_totals.index, title="Quantity", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```




The quantity feature significantly influences the target variable, total spent. When the total spent percentage is divided into five levels ranging from 1 to 5, it varies substantially—from as low as 6.3% up to 35.3%. This indicates that the quantity feature is a strong predictor of total spent.

that as the quantity increases, the total spent tends to rise noticeably, highlighting a strong positive relationship between quantity and total Spent.

In [815]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Price_Per_Unit', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Price_Per_Unit')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Price_Per_Unit')
plt.legend(wedges, location_totals.index, title="Price_Per_Unit", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```



The price per unit feature plays an important role in determining the total spent. When the total spent percentage is divided into five levels from 1 to 5, it varies widely—from 6.3% up to 35.3%. This variation shows that changes in the price per unit are closely associated with changes in total spending, indicating that higher unit prices generally lead to higher total expenditures.

In [816]:

```
nan_transaction_date_count = df['Transaction_Date'].isna().sum()
print(f"Number of rows where 'transaction_date' is NaN: {nan_transaction_date_count}")
```

Number of rows where 'transaction\_date' is NaN: 460

Out of 10,000 rows, 460 have a missing transaction\_date value, representing only 4.6% of the data. Since this is a relatively small proportion, these rows can be safely dropped without significantly impacting the overall analysis.

relatively small proportion, these rows can be safely dropped without significantly impacting the overall analysis.

In [817]:

```
df.dropna(subset=['Transaction_Date'], inplace=True)
```

In [818]:


```
df['Transaction_Date'] = pd.to_datetime(df['Transaction_Date'])
df['Day'] = df['Transaction_Date'].dt.day_name()
df['Month'] = df['Transaction_Date'].dt.month
```

To enhance the analysis, I plan to extract the day of the week and month from the transaction\_date feature. These time-based features are likely to be important, as customer behavior can vary depending on the day or month. For example, sales may be higher on weekends or during specific months. By creating these features, we can later visualize the data to verify whether such patterns exist.

In [819]:


```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Day', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Day')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Day')
plt.legend(wedges, location_totals.index, title="Day", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```

Total Spent by Day



In [820]:


```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Day', 'Total_Spent'], inplace=True)
month_totals = df_copy.groupby('Day')['Total_Spent'].sum()
plt.figure(figsize=(10, 6))
plt.plot(month_totals.index, month_totals.values, marker='o', linestyle='-', color='skyblue')
plt.title('Total Spent by Days')
plt.xlabel('Day')
plt.ylabel('Total Spent')
plt.xticks(month_totals.index)
plt.grid(True)
plt.show()
```




The total spent is not very affected by which day it is. But it is still considerable.

In [821]:


```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Month', 'Total_Spent'], inplace=True)
location_totals = df_copy.groupby('Month')['Total_Spent'].sum()
plt.figure(figsize=(8, 8))
wedges, texts, autotexts = plt.pie(location_totals, labels=location_totals.index, autopct='%1.1f%%', startangle=140, textprops={'color': 'black'})
plt.title('Total Spent by Month')
plt.legend(wedges, location_totals.index, title="Month", loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.show()
```





In [822]:

```
df_copy = df.copy()
df_copy['Total_Spent'] = pd.to_numeric(df_copy['Total_Spent'], errors='coerce')
df_copy.dropna(subset=['Month', 'Total_Spent'], inplace=True)
month_totals = df_copy.groupby('Month')['Total_Spent'].sum()
plt.figure(figsize=(10, 6))
plt.plot(month_totals.index, month_totals.values, marker='o', linestyle='-', color='skyblue')
plt.title('Total Spent by Month')
plt.xlabel('Month')
plt.ylabel('Total Spent')
plt.xticks(month_totals.index)
plt.grid(True)
plt.show()
```





The total spent is not very affected by which month we are in. But it is still considerable.

In [823]:

```
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 9540 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Transaction_ID  9540 non-null   object  
 1   Item              8613 non-null   object  
 2   Quantity          9086 non-null   object  
 3   Price_Per_Unit    9034 non-null   object  
 4   Total_Spent       9064 non-null   object  
 5   Payment_Method    6525 non-null   object  
 6   Location          5761 non-null   object  
 7   Transaction_Date  9540 non-null   datetime64[ns] 
 8   Day               9540 non-null   object  
 9   Month              9540 non-null   int32  
dtypes: datetime64[ns](1), int32(1), object(8)
memory usage: 782.6+ KB
```

In [824]:

```
df.drop(['Location', 'Payment_Method'], axis=1, inplace=True)

df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 9540 entries, 0 to 9999
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Transaction_ID  9540 non-null   object  
 1   Item              8613 non-null   object  
 2   Quantity          9086 non-null   object  
 3   Price_Per_Unit    9034 non-null   object  
 4   Total_Spent       9064 non-null   object  
 5   Transaction_Date  9540 non-null   datetime64[ns] 
 6   Day               9540 non-null   object  
 7   Month              9540 non-null   int32  
dtypes: datetime64[ns](1), int32(1), object(6)
memory usage: 633.5+ KB
```

The following is a code to check price per each item.

In [825]:

```
if 'Price_Per_Unit' in df.columns and 'Item' in df.columns:
    item_prices = df[['Item', 'Price_Per_Unit']].copy()
    item_prices.drop_duplicates(inplace=True)
    item_prices['Price_Per_Unit'] = pd.to_numeric(item_prices['Price_Per_Unit'], errors='coerce')
    item_prices.dropna(subset=['Price_Per_Unit'], inplace=True)
    print("Item Name and Price Per Unit:")
    for index, row in item_prices.iterrows():
        print(f"Item: {row['Item']}, Price Per Unit: {row['Price_Per_Unit']:.2f}")
else:
    print("The dataframe does not contain both 'Item' and 'Price_Per_Unit' columns.")
```

```
Item Name and Price Per Unit:
Item: Coffee, Price Per Unit: 2.00
Item: Cake, Price Per Unit: 3.00
Item: Cookie, Price Per Unit: 1.00
Item: Salad, Price Per Unit: 5.00
Item: Smoothie, Price Per Unit: 4.00
```

```
Item: nan, Price Per Unit: 3.00
Item: Sandwich, Price Per Unit: 4.00
Item: nan, Price Per Unit: 1.50
Item: Juice, Price Per Unit: 3.00
Item: nan, Price Per Unit: 2.00
Item: nan, Price Per Unit: 1.00
Item: Tea, Price Per Unit: 1.50
Item: nan, Price Per Unit: 5.00
Item: nan, Price Per Unit: 4.00
```

**Price Per Unit of Cookie is uniquely shared with one nan. Therefore we can use this info to fix the nan value that has Price Per Unit = 1. Similarly, we apply the same procedure for Tea and Salad.**

In [826]:

```
df.loc[(df['Price_Per_Unit'] == 1) & (df['Item'].isna()), 'Item'] = 'Cookie'
```

In [827]:

```
df.loc[(df['Price_Per_Unit'] == 1.5) & (df['Item'].isna()), 'Item'] = 'Tea'
```

In [828]:

```
df.loc[(df['Price_Per_Unit'] == 5) & (df['Item'].isna()), 'Item'] = 'Salad'
```

In [829]:

```
if 'Price_Per_Unit' in df.columns and 'Item' in df.columns:
    item_prices = df[['Item', 'Price_Per_Unit']].copy()
    item_prices['Price_Per_Unit'] = pd.to_numeric(item_prices['Price_Per_Unit'], errors='coerce')
    item_prices.dropna(subset=['Price_Per_Unit'], inplace=True)
    nan_item_prices = item_prices[item_prices['Item'].isna()]['Price_Per_Unit'].unique()
    for price in nan_item_prices:
        possible_items = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].unique()
        if len(possible_items) > 0:
            item_distribution = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].value_counts(normalize=True)
            nan_indices = df.index[(df['Item'].isna()) & (pd.to_numeric(df['Price_Per_Unit'], errors='coerce') == price)]
            if not item_distribution.empty and len(nan_indices) > 0:
                df.loc[nan_indices, 'Item'] = np.random.choice(item_distribution.index,
                                                               size=len(nan_indices),
                                                               p=item_distribution.values)
        else:
            print(f"Warning: No valid items found for Price_Per_Unit: {price}. NaNs for this price will remain.")
    print("Unique Item Name and Price Per Unit after filling NaNs:")
    item_prices_filled = df[['Item', 'Price_Per_Unit']].copy()
    item_prices_filled['Price_Per_Unit'] = pd.to_numeric(item_prices_filled['Price_Per_Unit'], errors='coerce')
    item_prices_filled.dropna(subset=['Item', 'Price_Per_Unit'], inplace=True)
    item_prices_filled['Item_lower'] = item_prices_filled['Item'].str.lower()
    unique_item_prices = item_prices_filled.drop_duplicates(subset=['Item_lower', 'Price_Per_Unit'])
    for index, row in unique_item_prices.iterrows():
        print(f"Item: {row['Item']}, Price Per Unit: {row['Price_Per_Unit']:.2f}")
else:
    print("The dataframe does not contain both 'Item' and 'Price_Per_Unit' columns.")
```

Unique Item Name and Price Per Unit after filling NaNs:

```
Item: Coffee, Price Per Unit: 2.00
Item: Cake, Price Per Unit: 3.00
Item: Cookie, Price Per Unit: 1.00
Item: Salad, Price Per Unit: 5.00
Item: Smoothie, Price Per Unit: 4.00
Item: Juice, Price Per Unit: 3.00
Item: Sandwich, Price Per Unit: 4.00
Item: Tea, Price Per Unit: 1.50
```

Now the items are filled successfully when one of price per unit or item is nan.

In the following we will deal with price per unit and item being both nan at the same time in an instance.

In [830]:

```
df[df['Item'].isna()]
```

Out[830]:

	Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Day	Month
118	TXN_4633784	NaN	5	NaN	15.0	2023-02-06	Monday	2
151	TXN_4031509	NaN	4	NaN	16.0	2023-01-04	Wednesday	1
289	TXN_3495950	NaN	4	NaN	6.0	2023-02-19	Sunday	2
334	TXN_2523298	NaN	4	NaN	6.0	2023-03-25	Saturday	3
550	TXN_4186681	NaN	4	NaN	6.0	2023-05-24	Wednesday	5
750	TXN_5787508	NaN	3	NaN	9.0	2023-07-23	Sunday	7
818	TXN_7940202	NaN	1	NaN	4.0	2023-07-23	Sunday	7
1154	TXN_2473090	NaN	2	NaN	3.0	2023-03-03	Friday	3
1337	TXN_5031214	NaN	5	NaN	5.0	2023-07-29	Saturday	7
1377	TXN_8396271	NaN	2	NaN	2.0	2023-09-12	Tuesday	9
1589	TXN_5245399	NaN	5	NaN	10.0	2023-12-25	Monday	12
1761	TXN_3611851	NaN	4	NaN	NaN	2023-02-09	Thursday	2
1786	TXN_1923349	NaN	4	NaN	6.0	2023-07-06	Thursday	7
2002	TXN_5206049	NaN	3	NaN	3.0	2023-06-24	Saturday	6
2227	TXN_3200203	NaN	2	NaN	8.0	2023-12-04	Monday	12
2289	TXN_7524977	NaN	4	NaN	NaN	2023-12-09	Saturday	12
2596	TXN_4844386	NaN	5	NaN	15.0	2023-10-28	Saturday	10
2610	TXN_8266689	NaN	5	NaN	15.0	2023-12-03	Sunday	12
2962	TXN_9702662	NaN	4	NaN	16.0	2023-07-10	Monday	7
3013	TXN_1842697	NaN	5	NaN	15.0	2023-10-25	Wednesday	10
3404	TXN_7797231	NaN	5	NaN	10.0	2023-05-05	Friday	5
3434	TXN_6457997	NaN	1	NaN	4.0	2023-12-12	Tuesday	12
3666	TXN_8616276	NaN	2	NaN	3.0	2023-07-22	Saturday	7
3739	TXN_4849180	NaN	5	NaN	15.0	2023-10-14	Saturday	10
3779	TXN_7376255	NaN	NaN	NaN	25.0	2023-05-27	Saturday	5
3900	TXN_5093855	NaN	4	NaN	6.0	2023-11-23	Thursday	11
4092	TXN_1840897	NaN	1	NaN	5.0	2023-06-03	Saturday	6
4152	TXN_9646000	NaN	2	NaN	NaN	2023-12-14	Thursday	12
4621	TXN_7844352	NaN	2	NaN	6.0	2023-07-26	Wednesday	7
5039	TXN_9514452	NaN	5	NaN	10.0	2023-07-29	Saturday	7
5891	TXN_9618962	NaN	3	NaN	6.0	2023-11-08	Wednesday	11
5991	TXN_2913107	NaN	4	NaN	8.0	2023-05-20	Saturday	5
6177	TXN_3232279	NaN	4	NaN	16.0	2023-05-30	Tuesday	5
6345	TXN_4208919	NaN	3	NaN	12.0	2023-05-30	Tuesday	5
6429	TXN_2536573	NaN	2	NaN	8.0	2023-06-24	Saturday	6
6849	TXN_7928378	NaN	3	NaN	12.0	2023-07-04	Tuesday	7
7478	TXN_9710982	NaN	4	NaN	16.0	2023-03-03	Friday	3

Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Wednesday	Month
8129	TXN_8384699	NaN	3	NaN	15.0	2023-12-16	Saturday 12
8458	TXN_7595907	NaN	5	NaN	15.0	2023-04-24	Monday 4
8463	TXN_2150872	NaN	5	NaN	15.0	2023-03-18	Saturday 3
8817	TXN_8365530	NaN	4	NaN	16.0	2023-03-10	Friday 3
8959	TXN_3803063	NaN	4	NaN	12.0	2023-11-23	Thursday 11
9174	TXN_5935353	NaN	2	NaN	8.0	2023-05-01	Monday 5
9234	TXN_8918253	NaN	3	NaN	12.0	2023-09-07	Thursday 9
9425	TXN_2065683	NaN	3	NaN	6.0	2023-09-21	Thursday 9
9673	TXN_2480808	NaN	1	NaN	4.0	2023-03-30	Thursday 3
9717	TXN_3334632	NaN	1	NaN	2.0	2023-11-20	Monday 11
9819	TXN_1208561	NaN	NaN	NaN	20.0	2023-08-19	Saturday 8
9820	TXN_8751702	NaN	5	NaN	15.0	2023-02-13	Monday 2
9996	TXN_9659401	NaN	3	NaN	3.0	2023-06-02	Friday 6

In the following, I want to replace the NaN of the price per unit by the value calculated. This value is calculated when dividing the total spent by the quantity.

In [831]:

```
df['Quantity'] = pd.to_numeric(df['Quantity'], errors='coerce')
df['Price_Per_Unit'] = pd.to_numeric(df['Price_Per_Unit'], errors='coerce')
df['Total_Spent'] = pd.to_numeric(df['Total_Spent'], errors='coerce')
df['Total_Spent'].fillna(df['Quantity'] * df['Price_Per_Unit'], inplace=True)
mask_fill_price_per_unit = df['Price_Per_Unit'].isna() & df['Total_Spent'].notna() & df['Quantity'].notna() & (df['Quantity'] != 0)
df.loc[mask_fill_price_per_unit, 'Price_Per_Unit'] = df.loc[mask_fill_price_per_unit, 'Total_Spent'] / df.loc[mask_fill_price_per_unit, 'Quantity']
mask_fill_quantity = df['Quantity'].isna() & df['Total_Spent'].notna() & df['Price_Per_Unit'].notna() & (df['Price_Per_Unit'] != 0)
df.loc[mask_fill_quantity, 'Quantity'] = df.loc[mask_fill_quantity, 'Total_Spent'] / df.loc[mask_fill_quantity, 'Price_Per_Unit']
df.head()
print(df.isna().sum())
```

```
Transaction_ID      0
Item                51
Quantity            36
Price_Per_Unit     35
Total_Spent         39
Transaction_Date   0
Day                 0
Month               0
dtype: int64
```

<ipython-input-831-2522293678>:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Total_Spent'].fillna(df['Quantity'] * df['Price_Per_Unit'], inplace=True)
```

In [832]:

```
df[df['Price_Per_Unit'].isna()]
```

Out[832]:

Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Day	Month	
Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Day	Month	
65	TXN_4987129	Sandwich	3.0	NaN	NaN	2023-10-20	Friday	10
629	TXN_9289174	Cake	NaN	NaN	12.0	2023-12-30	Saturday	12
912	TXN_1575608	Sandwich	NaN	NaN	20.0	2023-01-05	Thursday	1
1008	TXN_7225428	Tea	NaN	NaN	3.0	2023-03-07	Tuesday	3
1482	TXN_3593060	Smoothie	NaN	NaN	16.0	2023-03-05	Sunday	3
1674	TXN_9367492	Tea	2.0	NaN	NaN	2023-06-19	Monday	6
1761	TXN_3611851	NaN	4.0	NaN	NaN	2023-02-09	Thursday	2
2229	TXN_8498613	Sandwich	2.0	NaN	NaN	2023-11-08	Wednesday	11
2289	TXN_7524977	NaN	4.0	NaN	NaN	2023-12-09	Saturday	12
2330	TXN_3849488	Salad	NaN	NaN	5.0	2023-03-01	Wednesday	3
2585	TXN_1259340	Tea	3.0	NaN	NaN	2023-02-24	Friday	2
3162	TXN_3577949	Cake	3.0	NaN	NaN	2023-04-25	Tuesday	4
3598	TXN_2857444	Smoothie	1.0	NaN	NaN	2023-05-10	Wednesday	5
3635	TXN_6177081	Cookie	NaN	NaN	1.0	2023-07-26	Wednesday	7
3779	TXN_7376255	NaN	NaN	NaN	25.0	2023-05-27	Saturday	5
4021	TXN_6424202	Cookie	2.0	NaN	NaN	2023-11-20	Monday	11
4152	TXN_9646000	NaN	2.0	NaN	NaN	2023-12-14	Thursday	12
5639	TXN_6206792	Tea	NaN	NaN	6.0	2023-10-13	Friday	10
5845	TXN_8388462	Smoothie	NaN	NaN	8.0	2023-08-19	Saturday	8
6661	TXN_5308047	Cookie	NaN	NaN	1.0	2023-10-28	Saturday	10
6674	TXN_3071092	Coffee	NaN	NaN	4.0	2023-09-08	Friday	9
7035	TXN_8872984	Salad	5.0	NaN	NaN	2023-08-23	Wednesday	8
7230	TXN_5118799	Cookie	2.0	NaN	NaN	2023-04-23	Sunday	4
7244	TXN_2253622	Sandwich	5.0	NaN	NaN	2023-09-30	Saturday	9
7568	TXN_3705445	Cookie	5.0	NaN	NaN	2023-09-13	Wednesday	9
7597	TXN_1082717	NaN	NaN	NaN	9.0	2023-12-13	Wednesday	12
7686	TXN_6905143	Coffee	5.0	NaN	NaN	2023-08-09	Wednesday	8
8128	TXN_6105807	Smoothie	3.0	NaN	NaN	2023-01-18	Wednesday	1
8497	TXN_1525583	Sandwich	3.0	NaN	NaN	2023-05-20	Saturday	5
8669	TXN_7764304	Cake	NaN	NaN	3.0	2023-01-09	Monday	1
8975	TXN_8077367	Tea	NaN	NaN	7.5	2023-02-27	Monday	2
9212	TXN_7322317	Tea	1.0	NaN	NaN	2023-03-16	Thursday	3
9464	TXN_5744797	Cake	NaN	NaN	3.0	2023-12-20	Wednesday	12
9819	TXN_1208561	NaN	NaN	NaN	20.0	2023-08-19	Saturday	8
9893	TXN_3809533	Juice	2.0	NaN	NaN	2023-02-02	Thursday	2

All remaining instances have 2 NaN that can't be replaced by an exact value 100%. Therefore, it is better if we remove it but first we will try to replace NaN in item.

In [833]:

```
df.loc[(df['Price_Per_Unit'] == 5) & (df['Item'].isna()), 'Item'] = 'Salad'
df.loc[(df['Price_Per_Unit'] == 1.5) & (df['Item'].isna()), 'Item'] = 'Tea'
df.loc[(df['Price_Per_Unit'] == 1) & (df['Item'].isna()), 'Item'] = 'Cookie'
if 'Price_Per_Unit' in df.columns and 'Item' in df.columns:
    item_prices = df[['Item', 'Price_Per_Unit']].copy()
    item_prices['Price_Per_Unit'] = pd.to_numeric(item_prices['Price_Per_Unit'], errors='coerce')
    item_prices.dropna(subset=['Price_Per_Unit'], inplace=True)
```

```

nan_item_prices = item_prices[item_prices['Item'].isna()]['Price_Per_Unit'].unique()
for price in nan_item_prices:
    possible_items = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].unique()
    if len(possible_items) > 0:
        item_distribution = item_prices[(item_prices['Price_Per_Unit'] == price) & (item_prices['Item'].notna())]['Item'].value_counts(normalize=True)
        nan_indices = df.index[(df['Item'].isna()) & (pd.to_numeric(df['Price_Per_Unit'], errors='coerce') == price)]
        if not item_distribution.empty and len(nan_indices) > 0:
            df.loc[nan_indices, 'Item'] = np.random.choice(item_distribution.index, size=len(nan_indices), p=item_distribution.values)
    else:
        print(f"Warning: No valid items found for Price_Per_Unit: {price}. NaNs for this price will remain.")
print("Unique Item Name and Price Per Unit after filling NaNs:")
item_prices_filled = df[['Item', 'Price_Per_Unit']].copy()
item_prices_filled['Price_Per_Unit'] = pd.to_numeric(item_prices_filled['Price_Per_Unit'], errors='coerce')
item_prices_filled.dropna(subset=['Item', 'Price_Per_Unit'], inplace=True)
item_prices_filled['Item_lower'] = item_prices_filled['Item'].str.lower()
unique_item_prices = item_prices_filled.drop_duplicates(subset=['Item_lower', 'Price_Per_Unit'])

for index, row in unique_item_prices.iterrows():
    print(f"Item: {row['Item']}, Price Per Unit: {row['Price_Per_Unit']:.2f}")

else:
    print("The dataframe does not contain both 'Item' and 'Price_Per_Unit' columns.")

```

Unique Item Name and Price Per Unit after filling NaNs:

Item: Coffee, Price Per Unit: 2.00  
 Item: Cake, Price Per Unit: 3.00  
 Item: Cookie, Price Per Unit: 1.00  
 Item: Salad, Price Per Unit: 5.00  
 Item: Smoothie, Price Per Unit: 4.00  
 Item: Juice, Price Per Unit: 3.00  
 Item: Sandwich, Price Per Unit: 4.00  
 Item: Tea, Price Per Unit: 1.50

In [834]:

```
df[df['Item'].isna()]
```

Out[834]:

	Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Day	Month
1761	TXN_3611851	NaN	4.0	NaN	NaN	2023-02-09	Thursday	2
2289	TXN_7524977	NaN	4.0	NaN	NaN	2023-12-09	Saturday	12
3779	TXN_7376255	NaN	NaN	NaN	25.0	2023-05-27	Saturday	5
4152	TXN_9646000	NaN	2.0	NaN	NaN	2023-12-14	Thursday	12
7597	TXN_1082717	NaN	NaN	NaN	9.0	2023-12-13	Wednesday	12
9819	TXN_1208561	NaN	NaN	NaN	20.0	2023-08-19	Saturday	8

In [835]:

```

rows_with_nan = df[df.isna().any(axis=1)]
num_rows_with_nan = len(rows_with_nan)
print(f"Number of rows with at least one NaN value: {num_rows_with_nan}")
print("\nRows with NaN values:")
rows_with_nan

```

Number of rows with at least one NaN value: 55

Rows with NaN values:

Out [835]:

	Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Day	Month
65	TXN_4987129	Sandwich	3.0	NaN	NaN	2023-10-20	Friday	10
236	TXN_8562645	Salad	NaN	5.0	NaN	2023-05-18	Thursday	5
278	TXN_3229409	Juice	NaN	3.0	NaN	2023-04-15	Saturday	4
629	TXN_9289174	Cake	NaN	NaN	12.0	2023-12-30	Saturday	12
641	TXN_2962976	Juice	NaN	3.0	NaN	2023-03-17	Friday	3
738	TXN_8696094	Sandwich	NaN	4.0	NaN	2023-05-14	Sunday	5
912	TXN_1575608	Sandwich	NaN	NaN	20.0	2023-01-05	Thursday	1
1008	TXN_7225428	Tea	NaN	NaN	3.0	2023-03-07	Tuesday	3
1482	TXN_3593060	Smoothie	NaN	NaN	16.0	2023-03-05	Sunday	3
1674	TXN_9367492	Tea	2.0	NaN	NaN	2023-06-19	Monday	6
1761	TXN_3611851	NaN	4.0	NaN	NaN	2023-02-09	Thursday	2
2229	TXN_8498613	Sandwich	2.0	NaN	NaN	2023-11-08	Wednesday	11
2289	TXN_7524977	NaN	4.0	NaN	NaN	2023-12-09	Saturday	12
2330	TXN_3849488	Salad	NaN	NaN	5.0	2023-03-01	Wednesday	3
2585	TXN_1259340	Tea	3.0	NaN	NaN	2023-02-24	Friday	2
2796	TXN_9188692	Cake	NaN	3.0	NaN	2023-12-01	Friday	12
3162	TXN_3577949	Cake	3.0	NaN	NaN	2023-04-25	Tuesday	4
3203	TXN_4565754	Smoothie	NaN	4.0	NaN	2023-10-06	Friday	10
3224	TXN_6297232	Coffee	NaN	2.0	NaN	2023-04-07	Friday	4
3401	TXN_3251829	Tea	NaN	1.5	NaN	2023-07-25	Tuesday	7
3598	TXN_2857444	Smoothie	1.0	NaN	NaN	2023-05-10	Wednesday	5
3635	TXN_6177081	Cookie	NaN	NaN	1.0	2023-07-26	Wednesday	7
3779	TXN_7376255	NaN	NaN	NaN	25.0	2023-05-27	Saturday	5
4021	TXN_6424202	Cookie	2.0	NaN	NaN	2023-11-20	Monday	11
4152	TXN_9646000	NaN	2.0	NaN	NaN	2023-12-14	Thursday	12
4257	TXN_6470865	Coffee	NaN	2.0	NaN	2023-09-18	Monday	9
5639	TXN_6206792	Tea	NaN	NaN	6.0	2023-10-13	Friday	10
5841	TXN_5884081	Cookie	NaN	1.0	NaN	2023-07-05	Wednesday	7
5845	TXN_8388462	Smoothie	NaN	NaN	8.0	2023-08-19	Saturday	8
6661	TXN_5308047	Cookie	NaN	NaN	1.0	2023-10-28	Saturday	10
6674	TXN_3071092	Coffee	NaN	NaN	4.0	2023-09-08	Friday	9
7029	TXN_4628338	Coffee	NaN	2.0	NaN	2023-12-25	Monday	12
7035	TXN_8872984	Salad	5.0	NaN	NaN	2023-08-23	Wednesday	8
7230	TXN_5118799	Cookie	2.0	NaN	NaN	2023-04-23	Sunday	4
7244	TXN_2253622	Sandwich	5.0	NaN	NaN	2023-09-30	Saturday	9
7297	TXN_9944500	Smoothie	NaN	4.0	NaN	2023-01-03	Tuesday	1
7568	TXN_3705445	Cookie	5.0	NaN	NaN	2023-09-13	Wednesday	9
7597	TXN_1082717	NaN	NaN	NaN	9.0	2023-12-13	Wednesday	12
7686	TXN_6905143	Coffee	5.0	NaN	NaN	2023-08-09	Wednesday	8
8021	TXN_2428781	Salad	NaN	5.0	NaN	2023-05-09	Tuesday	5
8128	TXN_6105807	Smoothie	3.0	NaN	NaN	2023-01-18	Wednesday	1
8443	TXN_2023651	Sandwich	NaN	4.0	NaN	2023-05-25	Thursday	5
8465	TXN_9669616	Coffee	NaN	2.0	NaN	2023-06-03	Saturday	6

8479	TXN_1547215	Sandwich	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Monday	Month
8497	TXN_1525583	Sandwich	3.0	NaN	NaN	2023-05-20	Saturday	5
8574	TXN_2546684	Juice	NaN	3.0	NaN	2023-04-08	Saturday	4
8669	TXN_7764304	Cake	NaN	NaN	3.0	2023-01-09	Monday	1
8732	TXN_4550558	Cookie	NaN	1.0	NaN	2023-08-04	Friday	8
8975	TXN_8077367	Tea	NaN	NaN	7.5	2023-02-27	Monday	2
9212	TXN_7322317	Tea	1.0	NaN	NaN	2023-03-16	Thursday	3
9464	TXN_5744797	Cake	NaN	NaN	3.0	2023-12-20	Wednesday	12
9590	TXN_9924732	Sandwich	NaN	4.0	NaN	2023-01-18	Wednesday	1
9819	TXN_1208561	NaN	NaN	NaN	20.0	2023-08-19	Saturday	8
9869	TXN_1975184	Coffee	NaN	2.0	NaN	2023-01-15	Sunday	1
9893	TXN_3809533	Juice	2.0	NaN	NaN	2023-02-02	Thursday	2

In [836]:

```
mask_price_nan_item_not_nan = df['Price_Per_Unit'].isna() & df['Item'].notna()
items_with_price_nan = df.loc[mask_price_nan_item_not_nan, 'Item'].unique()
item_known_prices = {}
for item in items_with_price_nan:
    item_valid_prices = df.loc[
        (df['Item'] == item) & df['Price_Per_Unit'].notna(),
        'Price_Per_Unit'
    ].astype(float).unique()
    if len(item_valid_prices) > 0:
        item_known_prices[item] = item_valid_prices[0]
for index, row in df.loc[mask_price_nan_item_not_nan].iterrows():
    item = row['Item']
    if item in item_known_prices:
        df.loc[index, 'Price_Per_Unit'] = item_known_prices[item]

df[df['Price_Per_Unit'].isna()]
```

Out[836]:

	Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	Transaction_Date	Day	Month
1761	TXN_3611851	NaN	4.0	NaN	NaN	2023-02-09	Thursday	2
2289	TXN_7524977	NaN	4.0	NaN	NaN	2023-12-09	Saturday	12
3779	TXN_7376255	NaN	NaN	NaN	25.0	2023-05-27	Saturday	5
4152	TXN_9646000	NaN	2.0	NaN	NaN	2023-12-14	Thursday	12
7597	TXN_1082717	NaN	NaN	NaN	9.0	2023-12-13	Wednesday	12
9819	TXN_1208561	NaN	NaN	NaN	20.0	2023-08-19	Saturday	8

In [837]:

```
df.dropna(subset=['Item'], inplace=True)
df[df['Total_Spent'].isna()]
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9534 entries, 0 to 9999
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Transaction_ID  9534 non-null   object 
 1   Item              9534 non-null   object 
 2   Quantity          9501 non-null   float64
 3   Price_Per_Unit   9534 non-null   float64
 4   Total_Spent       9498 non-null   float64
 5   Transaction_Date 9534 non-null   datetime64[ns]
 6   Day               9534 non-null   object 
 7   Month              9534 non-null   int32
```

```
dtypes: datetime64[ns](1), float64(3), int32(1), object(3)
memory usage: 633.1+ KB
```

In [838]:

```
item_known_prices = df.dropna(subset=['Item', 'Price_Per_Unit']).groupby('Item')['Price_Per_Unit'].first().to_dict()
mask_fill_price_per_unit_by_item = df['Price_Per_Unit'].isna() & df['Item'].notna()
for index, row in df.loc[mask_fill_price_per_unit_by_item].iterrows():
    item = row['Item']
    if item in item_known_prices:
        df.loc[index, 'Price_Per_Unit'] = item_known_prices[item]
print("\nRows with NaN Price_Per_Unit after filling based on Item:")
print(df[df['Price_Per_Unit'].isna()])
rows_with_nan_after_item_fill = df[df.isna().any(axis=1)]
num_rows_with_nan_after_item_fill = len(rows_with_nan_after_item_fill)
print(f"\nNumber of rows with at least one NaN value after filling price by item: {num_rows_with_nan_after_item_fill}")
```

Rows with NaN Price\_Per\_Unit after filling based on Item:

Empty DataFrame

Columns: [Transaction\_ID, Item, Quantity, Price\_Per\_Unit, Total\_Spent, Transaction\_Date, Da

y, Month]

Index: []

Number of rows with at least one NaN value after filling price by item: 49

In [839]:

```
df['Quantity'] = pd.to_numeric(df['Quantity'], errors='coerce')
df['Price_Per_Unit'] = pd.to_numeric(df['Price_Per_Unit'], errors='coerce')
df['Total_Spent'] = pd.to_numeric(df['Total_Spent'], errors='coerce')
df['Total_Spent'].fillna(df['Quantity'] * df['Price_Per_Unit'], inplace=True)
mask_fill_price_per_unit = df['Price_Per_Unit'].isna() & df['Total_Spent'].notna() & df['Quantity'].notna() & (df['Quantity'] != 0)
df.loc[mask_fill_price_per_unit, 'Price_Per_Unit'] = df.loc[mask_fill_price_per_unit, 'Total_Spent'] / df.loc[mask_fill_price_per_unit, 'Quantity']
mask_fill_quantity = df['Quantity'].isna() & df['Total_Spent'].notna() & df['Price_Per_Unit'].notna() & (df['Price_Per_Unit'] != 0)
df.loc[mask_fill_quantity, 'Quantity'] = df.loc[mask_fill_quantity, 'Total_Spent'] / df.loc[mask_fill_quantity, 'Price_Per_Unit']
df.head()
df.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 9534 entries, 0 to 9999

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
---	---	-----	----
0	Transaction_ID	9534	non-null
1	Item	9534	non-null
2	Quantity	9514	non-null
3	Price_Per_Unit	9534	non-null
4	Total_Spent	9514	non-null
5	Transaction_Date	9534	non-null
6	Day	9534	non-null
7	Month	9534	non-null

dtypes: datetime64[ns](1), float64(3), int32(1), object(3)

memory usage: 633.1+ KB

```
<ipython-input-839-769373654>:4: FutureWarning: A value is trying to be set on a copy of a
DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the
intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Total_Spent'].fillna(df['Quantity'] * df['Price_Per_Unit'], inplace=True)
```

In [840]:

```

df.dropna(inplace=True)
print("DataFrame Info:")
df.info()
print("\nNaN values per column:")
print(df.isnull().sum())

```

DataFrame Info:

<class 'pandas.core.frame.DataFrame'>

Index: 9514 entries, 0 to 9999

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Transaction_ID	9514 non-null	object
1	Item	9514 non-null	object
2	Quantity	9514 non-null	float64
3	Price_Per_Unit	9514 non-null	float64
4	Total_Spent	9514 non-null	float64
5	Transaction_Date	9514 non-null	datetime64[ns]
6	Day	9514 non-null	object
7	Month	9514 non-null	int32

dtypes: datetime64[ns](1), float64(3), int32(1), object(3)

memory usage: 631.8+ KB

NaN values per column:

Transaction_ID	0
Item	0
Quantity	0
Price_Per_Unit	0
Total_Spent	0
Transaction_Date	0
Day	0
Month	0

dtype: int64

## EDA after data preprocessing

In [841]:

```

print("\nDataFrame Description:")
print(df.describe(include='all'))

```

DataFrame Description:

	Transaction_ID	Item	Quantity	Price_Per_Unit	Total_Spent	\
count	9514	9514	9514.000000	9514.000000	9514.000000	
unique	9514	8	NaN	NaN	NaN	
top	TXN_6170729	Coffee	NaN	NaN	NaN	
freq	1	1241	NaN	NaN	NaN	
mean	NaN	NaN	3.021232	2.948550	8.920380	
min	NaN	NaN	1.000000	1.000000	1.000000	
25%	NaN	NaN	2.000000	2.000000	4.000000	
50%	NaN	NaN	3.000000	3.000000	8.000000	
75%	NaN	NaN	4.000000	4.000000	12.000000	
max	NaN	NaN	5.000000	5.000000	25.000000	
std	NaN	NaN	1.420285	1.279259	6.003809	

	Transaction_Date	Day	Month
count	9514	9514	9514.000000
unique	NaN	7	NaN
top	NaN	Friday	NaN
freq	NaN	1383	NaN
mean	2023-07-01 23:14:35.593861632	NaN	6.524070
min	2023-01-01 00:00:00	NaN	1.000000
25%	2023-04-01 00:00:00	NaN	4.000000
50%	2023-07-02 00:00:00	NaN	7.000000
75%	2023-10-02 00:00:00	NaN	10.000000
max	2023-12-31 00:00:00	NaN	12.000000
std	NaN	NaN	3.449063

In [842]:

```
item_total_spent = df.groupby('Item')['Total_Spent'].sum().sort_values(ascending=False)
```

```


plt.figure(figsize=(12, 6))
sns.barplot(x=item_total_spent.index, y=item_total_spent.values)
plt.xticks(rotation=90)
plt.title('Total Spent by Item')
plt.xlabel('Item')
plt.ylabel('Total Spent')
plt.tight_layout()
plt.show()



plt.figure(figsize=(8, 6))
sns.scatterplot(x=df['Quantity'], y=df['Total_Spent'])
plt.title('Total Spent vs. Quantity')
plt.xlabel('Quantity')
plt.ylabel('Total Spent')
plt.show()

month_total_spent = df.groupby('Month')['Total_Spent'].sum()
plt.figure(figsize=(10, 6))
sns.lineplot(x=month_total_spent.index, y=month_total_spent.values)
plt.title('Total Spent by Month')
plt.xlabel('Month')
plt.ylabel('Total Spent')
plt.xticks(month_total_spent.index, ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']) # Label months
plt.grid(True)
plt.show()

df['Day_Name'] = df['Transaction_Date'].dt.day_name()
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
day_name_total_spent = df.groupby('Day_Name')['Total_Spent'].sum().reindex(day_order)
plt.figure(figsize=(10, 6))
sns.lineplot(x=day_name_total_spent.index, y=day_name_total_spent.values)
plt.title('Total Spent by Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Total Spent')
plt.grid(True)
plt.tight_layout()
plt.show()

```








In [843]:


```
numerical_cols = ['Quantity', 'Price_Per_Unit', 'Day_Name', 'Month']


for col in numerical_cols:
    if col in df.columns:
        plt.figure(figsize=(8, 6))
        sns.boxplot(x=df[col])
        plt.title(f'Box Plot of {col}')
        plt.xlabel(col)
        plt.show()
```

Box Plot of Quantity




Box Plot of Price\_Per\_Unit






Box Plot of Day\_Name



Box Plot of Month







## Dataset Balance

Since the dataset lacks timestamps (no hours), we cannot analyze hourly trends (so “peak hours” analysis is skipped). Therefore, we will plot our target (total spent) that will be predicted later on.

Since this is regression, we inspect *balance* by plotting a histogram (or KDE) of Total\_Spent. A highly skewed target (heavy outliers) would suggest careful metric choice.

In [844]:

```
y = df['Total_Spent'] # target variable
# histogram of Total_Spent
plt.figure(figsize=(6, 4))
sns.histplot(y, bins=15, kde=True)
plt.title('Distribution of Total Spent (Target)')
plt.xlabel('Total Spent ($)')
plt.ylabel('Frequency')
plt.show()
```



The range is between 1 and 25.

The median is around 7.5 and mean about 8.8.

It is clear that we are dealing with *positive-to-zero* skewness --> our target is nicely distributed around a central value. Therefore standard regression metrics must be used (MAE,MSE,RMSE,R<sup>2</sup>).

## Regression Metrics

1. Because this is a regression problem (continuous target), we will focus on regression metrics like MAE, MSE, RMSE, and R<sup>2</sup>.
2. Confusion matrix is not applicable here, as confusion matrix summarizes classification errors, not continuous regression predictions.
3. RMSE is in the same units as sales (\$) and penalizes large errors more than MAE does.

3. RMSE is in the same units as sales (\$) and penalizes large errors more than MAE does.
4. R<sup>2</sup> of 1.0 means perfect fit, whereas 0 means no predictive power beyond the mean.
5. Quantile (pinball) loss is not about where most of your data live—it's about which “slice” of the error distribution you care to optimize. In my case, I decided to optimize big spenders, so I choose  $\alpha=0.9$ .

$$L_\alpha(y, \hat{y}) = \begin{cases} \alpha(y - \hat{y}), & \text{if } y \geq \hat{y}, \\ (1 - \alpha)(\hat{y} - y), & \text{if } y < \hat{y} \end{cases}$$

## Standard Train-Test Split

Setting a random seed ensures that anyone running this code will get the same split and results.

In [845]:

```
from sklearn.model_selection import train_test_split

# these columns could be easily removed because I already extracted Day and Month columns
df.drop(columns=['Transaction_Date', 'Transaction_ID'], inplace=True)

# categorical variables (Item, Day) are one-hot encoded because computers understand numbers
# better than string (it will be used in Feature selection)
# used one-hot encoding to prevent association.
X = pd.get_dummies(df.drop(columns=['Total_Spent'])), columns=['Item', 'Day'], drop_first=True
X.drop(columns=['Day_Name'], inplace=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.4, random_state=42
)

# printing the shapes is used to confirm the size of each
print("Train set:", X_train.shape, "Test set:", X_test.shape)
```

Train set: (5708, 16) Test set: (3806, 16)

In [846]:

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9514 entries, 0 to 9999
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Quantity         9514 non-null   float64 
 1   Price_Per_Unit  9514 non-null   float64 
 2   Month            9514 non-null   int32  
 3   Item_Coffee      9514 non-null   bool   
 4   Item_Cookie      9514 non-null   bool   
 5   Item_Juice       9514 non-null   bool   
 6   Item_Salad       9514 non-null   bool   
 7   Item_Sandwich    9514 non-null   bool   
 8   Item_Smoothie    9514 non-null   bool   
 9   Item_Tea          9514 non-null   bool   
 10  Day_Monday       9514 non-null   bool   
 11  Day_Saturday     9514 non-null   bool   
 12  Day_Sunday        9514 non-null   bool   
 13  Day_Thursday     9514 non-null   bool   
 14  Day_Tuesday      9514 non-null   bool   
 15  Day_Wednesday    9514 non-null   bool   
dtypes: bool(13), float64(2), int32(1)
memory usage: 380.9 KB
```

## Feature Selection

In the following we perform feature selection to reduce dimensionality:

1. SelectKBest picks the top k features by a statistical test.
2. RFE (Recursive Feature Elimination) uses a model feature importances to iteratively drop the least important features.

features.

### 3. Sequential Forward/Backward Selection would add/remove features one by one.

I will not try Sequential Selection because it can be very slow.

In [847]:

```
from sklearn.feature_selection import SelectKBest, f_regression
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

# use SelectKBest with f_regression to select top 5 features
skb = SelectKBest(score_func=f_regression, k=5)
X_train_skb = skb.fit_transform(X_train, y_train)
X_test_skb = skb.transform(X_test)
print("After SelectKBest, Train shape:", X_train_skb.shape, "Test shape:", X_test_skb.shape)

# use RFE with a linear model to select top 5 features
rfe = RFE(estimator=LinearRegression(), n_features_to_select=5)
X_train_rfe = rfe.fit_transform(X_train, y_train)
X_test_rfe = rfe.transform(X_test)
print("After RFE, Train shape:", X_train_rfe.shape, "Test shape:", X_test_rfe.shape)
```

After SelectKBest, Train shape: (5708, 5) Test shape: (3806, 5)  
After RFE, Train shape: (5708, 5) Test shape: (3806, 5)

## Model Building and Evaluation/Comparison

Each model is trained on the same train set and evaluated on the same test set. We record MAE, MSE, RMSE, and R<sup>2</sup> for comparison.

In [848]:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, mean_pinball_loss
from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import GradientBoostingRegressor

models = {
    "Linear Regression": LinearRegression(),
    "MLP Regressor": MLPRegressor(hidden_layer_sizes=(200,), max_iter=300, random_state=42),
    "KNN Regressor": KNeighborsRegressor(n_neighbors=5),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

alpha = 0.9 # 90th percentile pinball loss

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
    r2 = r2_score(y_test, preds)
    qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

    results.append([name, mae, mse, rmse, r2, qloss])

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))
```

	Model	MAE	MSE	RMSE	R2	\
0	Linear Regression	1.292841	3.260440	1.805669	0.909003	
1	MLP Regressor	0.038853	0.002936	0.054181	0.999918	
2	KNN Regressor	0.451104	0.497501	0.705338	0.986115	

```

3 Gradient Boosting  0.031999  0.002167  0.046555  0.999940
PinballLoss(q=0.9)
0          0.633414
1          0.024154
2          0.229514
3          0.016485

```

From this, we see Gradient Boosting achieve very high accuracy ( $R^2$  near 1.0, extremely low error). This is because with the product relationship in the data, a tree can perfectly learn the rule.

One hot encoding tried to reduce the product relationship in the data, but it was not enough (mainly because of how the dataset was given initially).

Linear regression is less flexible ( $R^2 \approx 0.90$ ). It underperformed because relationships between data are not linear.

The KNN regressor also does well by locally interpolating neighbors. Prediction time increases rapidly with increased data size. Fortunately this dataset was not big.

## Hyperparameter tuning

In [849]:

```

from sklearn.model_selection import GridSearchCV
"""

from sklearn.linear_model import Ridge
ridge = Ridge()
ridge_params = {'alpha': [0.01, 0.1, 1, 10, 100, 1000]}
ridge_grid = GridSearchCV(ridge, ridge_params, cv=5, scoring='r2')
ridge_grid.fit(X_train, y_train)
print("Best Ridge Params:", ridge_grid.best_params_)
print("")

from sklearn.linear_model import Lasso
lasso = Lasso(max_iter=10000)
lasso_params = {'alpha': [0.001, 0.01, 0.1, 1, 10, 100, 1000]}
lasso_grid = GridSearchCV(lasso, lasso_params, cv=5, scoring='r2')
lasso_grid.fit(X_train, y_train)
print("Best Lasso Params:", lasso_grid.best_params_)

"""

from sklearn.neural_network import MLPRegressor
mlp = MLPRegressor(max_iter=1000, random_state=42)
mlp_params = {
    'hidden_layer_sizes': [(50,), (100,), (50,50)],
    'activation': ['relu', 'tanh'],
    'solver': ['adam', 'sgd'],
    'alpha': [0.0001, 0.001], # regularization
}
mlp_grid = GridSearchCV(mlp, mlp_params, cv=5, scoring='r2', n_jobs=-1)
mlp_grid.fit(X_train, y_train)
print("Best MLP Params:", mlp_grid.best_params_)
print("")

from sklearn.neighbors import KNeighborsRegressor
knn = KNeighborsRegressor()
knn_params = {
    'n_neighbors': [3, 5, 7, 10],
    'weights': ['uniform', 'distance'],
    'p': [1, 2] # 1: Manhattan, 2: Euclidean
}
knn_grid = GridSearchCV(knn, knn_params, cv=5, scoring='r2')
knn_grid.fit(X_train, y_train)
print("Best KNN Params:", knn_grid.best_params_)

```

Best Lasso Params: {'alpha': 0.01}

Best KNN Params: {'n\_neighbors': 10, 'p': 2, 'weights': 'distance'}

In [850]:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, mean_pinball_loss
from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import GradientBoostingRegressor

models = {
    "Linear Regression": LinearRegression(),
    "Lasso Regression": Lasso(alpha=0.01),
    "KNN Regressor": KNeighborsRegressor(n_neighbors=5),
    "KNN RegressorHT": KNeighborsRegressor(n_neighbors=10, weights='distance', p=2),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

alpha = 0.9 # 90th percentile pinball loss

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
    r2 = r2_score(y_test, preds)
    qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

    results.append([name, mae, mse, rmse, r2, qloss])

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))
```

	Model	MAE	MSE	RMSE	R2	\
0	Linear Regression	1.292841	3.260440	1.805669	0.909003	
1	Lasso Regression	1.287621	3.251464	1.803182	0.909254	
2	KNN Regressor	0.451104	0.497501	0.705338	0.986115	
3	KNN RegressorHT	0.097073	0.091894	0.303140	0.997435	
4	Gradient Boosting	0.031999	0.002167	0.046555	0.999940	

	PinballLoss(q=0.9)
0	0.633414
1	0.630333
2	0.229514
3	0.045008
4	0.016485

In [851]:

```
from sklearn.linear_model import Lasso
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform, randint

lasso_params = {
    'alpha': uniform(0.0001, 1) # Try small to moderate penalties
}
lasso_search = RandomizedSearchCV(
    Lasso(max_iter=10000), # Ensure enough iterations for convergence
    param_distributions=lasso_params,
    n_iter=20,
    cv=5,
    scoring='r2',
    random_state=42
)
lasso_search.fit(X_train, y_train)
print("Lasso Best Params:", lasso_search.best_params_)
```

```

from sklearn.neighbors import KNeighborsRegressor
knn_params = {
    'n_neighbors': randint(3, 20),
    'weights': ['uniform', 'distance'],
    'p': [1, 2]
}
knn_search = RandomizedSearchCV(
    KNeighborsRegressor(),
    knn_params,
    n_iter=20,
    cv=5,
    scoring='r2',
    random_state=42
)
knn_search.fit(X_train, y_train)
print("KNN Best Params:", knn_search.best_params_)

```

Lasso Best Params: {'alpha': np.float64(0.020684494295802446)}  
KNN Best Params: {'n\_neighbors': 16, 'p': 2, 'weights': 'distance'}

In [852]:

```

models = {
    "Linear Regression": LinearRegression(),
    "Lasso Regression": Lasso(alpha=0.02),
    "KNN Regressor": KNeighborsRegressor(n_neighbors=5),
    "KNN RegressorHT": KNeighborsRegressor(n_neighbors=16, weights='distance', p=2),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

alpha = 0.9 # 90th percentile pinball loss

results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    mae = mean_absolute_error(y_test, preds)
    mse = mean_squared_error(y_test, preds)
    rmse = mse**0.5
    r2 = r2_score(y_test, preds)
    qloss = mean_pinball_loss(y_test, preds, alpha=alpha)

    results.append([name, mae, mse, rmse, r2, qloss])

print(pd.DataFrame(results, columns=["Model", "MAE", "MSE", "RMSE", "R2", f"PinballLoss(q={alpha})"]))

```

	Model	MAE	MSE	RMSE	R2	\
0	Linear Regression	1.292841	3.260440	1.805669	0.909003	
1	Lasso Regression	1.290072	3.253756	1.803817	0.909190	
2	KNN Regressor	0.451104	0.497501	0.705338	0.986115	
3	KNN RegressorHT	0.089367	0.078033	0.279345	0.997822	
4	Gradient Boosting	0.031999	0.002167	0.046555	0.999940	
	PinballLoss(q=0.9)					
0		0.633414				
1		0.631718				
2		0.229514				
3		0.044664				
4		0.016485				

The results showed that Random Search outperformed Grid Search for the KNN Regressor, while Grid Search provided better results for Lasso Regression.

KNN involves multiple hyperparameters and less sensitivity to each one (suitable for Random Search which explores large spaces without testing all possibilities/combinations).

## Conclusion

We implemented a full ML pipeline for predicting cafe sales. Data cleaning resolved missing and inconsistent values, EDA guided our understanding (no peak-hour analysis possible without time data), and feature engineering added useful temporal categories. After splitting the data (reproducibly using a fixed random\_state geeksforgeeks.org ), we compared multiple regressors. Gradient Boosting and the neural net achieved almost perfect fit, highlighting that complex models captured the sales formula very well. We emphasized appropriate regression metrics (MAE, RMSE, R<sup>2</sup>, quantile loss). The code and visualizations above can be run in Google Colab or any Python environment for full exploration.